

Data Visualization with R Tidyverse tutorial

Ariane Ducellier

University of Washington - Fall 2025

What is tidyverse?

A collection of R packages designed for data science.

Basic packages:

- **ggplot2**: graphics
- **dplyr**: data manipulation
- **tidyr**: getting to tidy data
- **readr**: reading rectangular data (e.g. csv, tsv, fwf)
- **purrr**: working with functions and vectors
- **tibble**: a modern re-imagining of the data frame
- **stringr**: working with strings
- **forcats**: working with R factors to handle categorical variables

Additional packages associated to tidyverse need to be installed and loaded separately to import data, wrangle data, program and model.

Main concepts of data wrangling

- Understand.
- Format → Produce tidy data:
 - Every column is a variable.
 - Every row is an observation.
 - Every cell is a single value.
- Clean.
- Enrich.
- Validate.
- Analysis / Model → In our case, we are going to produce visuals to communicate information on the dataset to the viewer.

Benefits of data wrangling

- Organized and easily understandable data.
- Faster results.
- Better data flow for modeling or data visualization.
- Easier aggregation for insight extraction.
- Data quality.
- Data enriching.

Frameworks in Data Science: KDD

Knowledge Discovery in Databases

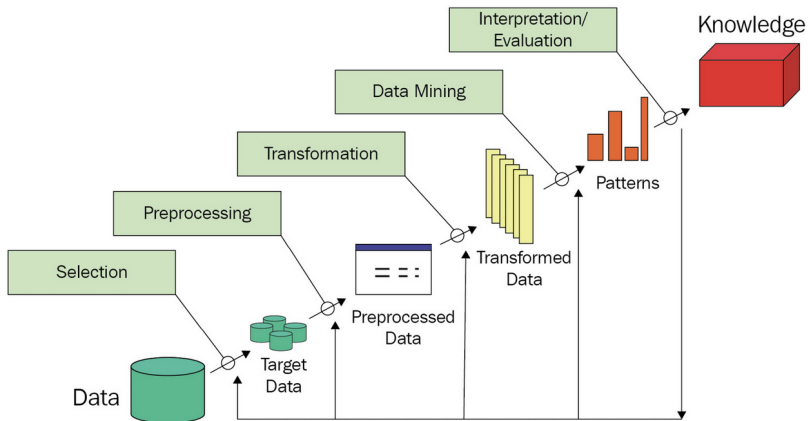


Figure from G.R. Santos, Data Wrangling with R.

Knowledge Discovery in Databases

- Getting the data.
- Selecting a subset of samples / variables of interest.
- Preprocessing (remove outliers, handle missing or noisy data).
- Transformation and formatting.
- Data mining (e.g. classification, clustering).
- Interpretation and evaluation.

Frameworks in Data Science: SEMMA

Sample, Explore, Modify, Model, and Assess

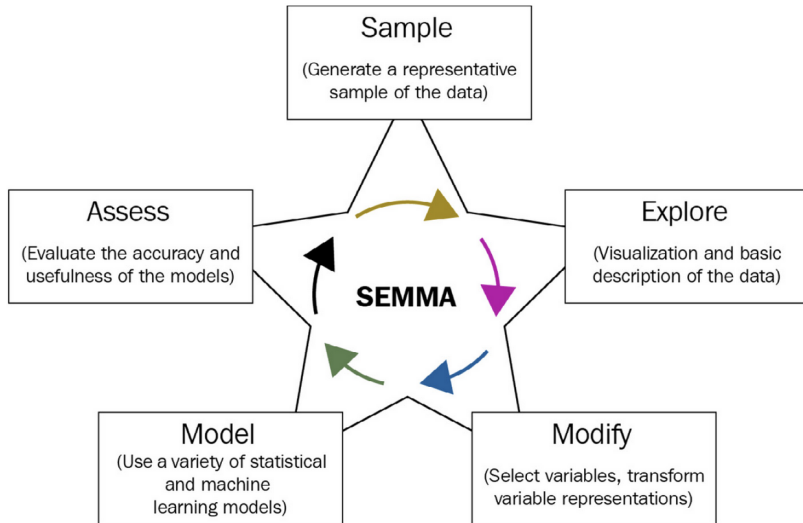


Figure from G.R. Santos, Data Wrangling with R.

Sample, Explore, Modify, Model, and Assess

Cyclic process:

- Sample (representative sample, but easy to work with).
- Explore (understand, visualize, describe, patterns and anomalies).
- Modify (data wrangling).
- Model (algorithms for predictions or insights on the data).
- Access (evaluate the results).

Frameworks in Data Science: CRISP-DM

Cross-Industry Standard Process for Data Mining




Figure from G.R. Santos, Data Wrangling with R.

Cross-Industry Standard Process for Data Mining

- Business understanding: Understand the problem and the business rules and specificities.
- Data understanding: Explore the data, find errors and missing data to assess quality.
- Data preparation: Data wrangling.
- Modeling: Analysis of the processed data.
- Evaluation: Assess whether the solution is aligned with the business requirements.
- Deployment: The model reaches its purpose.

Tibbles versus Data frames

- Tibbles do not change input variable types by default.
- Tibbles can have lists as columns.
- Tibbles can have non-standard variable names.
- Tibbles return another Tibble when slicing (and not a vector).

The pipe operator

The `magrittr` package provides the `%>%` operator as a shortcut for modifying an object in place:

```
df_iris <- iris %>%  
  group_by(Species) %>%  
  summarize_if(is.numeric, mean) %>%  
  ungroup() %>%  
  gather(measure, value, -Species) %>%  
  arrange(value)
```

=

```
df_iris <- group_by(iris, Species)  
df_iris <- summarize_if(df_iris, is.numeric, mean)  
df_iris <- ungroup(df_iris)  
df_iris <- gather(df_iris, measure, value, -Species)  
df_iris <- arrange(df_iris, value)
```