

# Ggplot2 tutorial - Command lines

Ariane Ducellier


7/18/2023

```
library(ggplot2)
library(gridExtra)
library(Lock5Data)
```

## Part 1 - Basic plotting in ggplot2

### Histograms

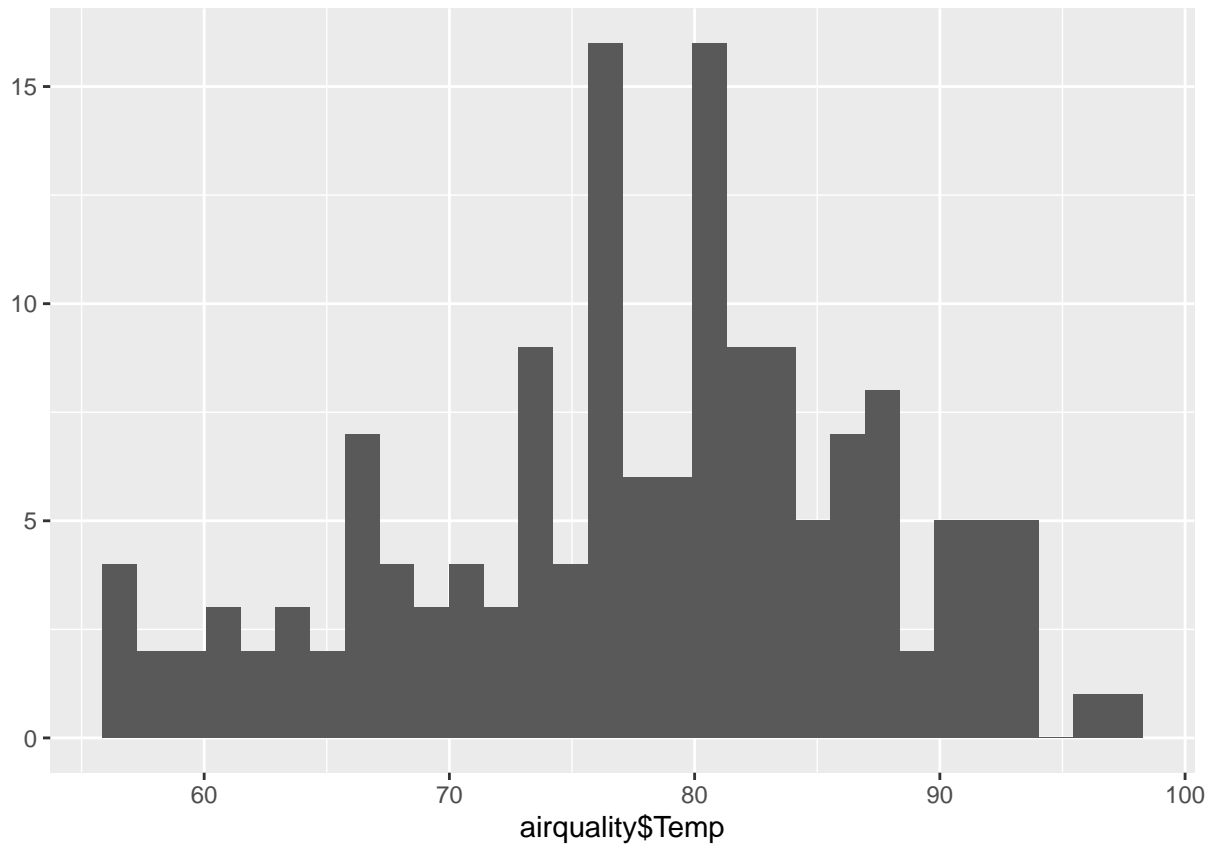
```
hist(airquality$Temp)
```



```
qplot(airquality$Temp)
```

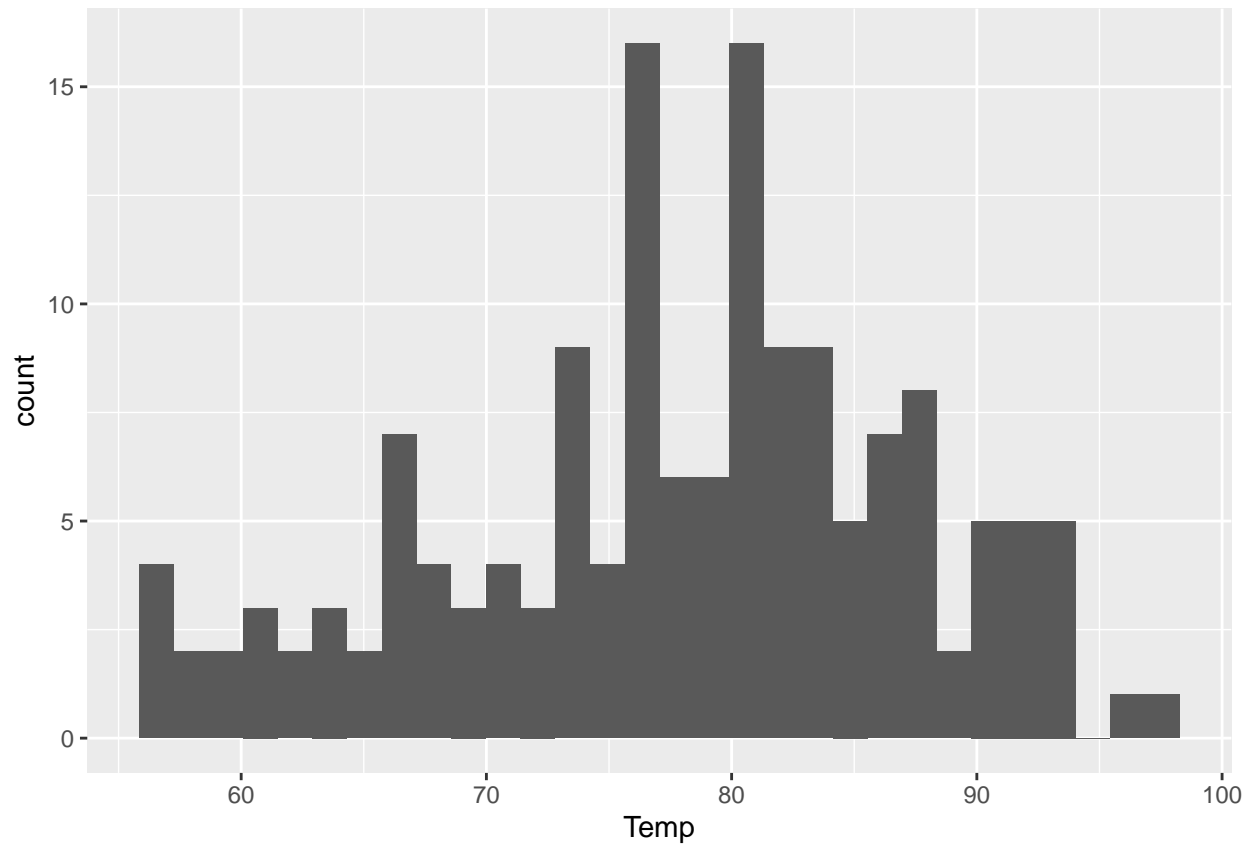
```
## Warning: 'qplot()' was deprecated in ggplot2 3.4.0.
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```



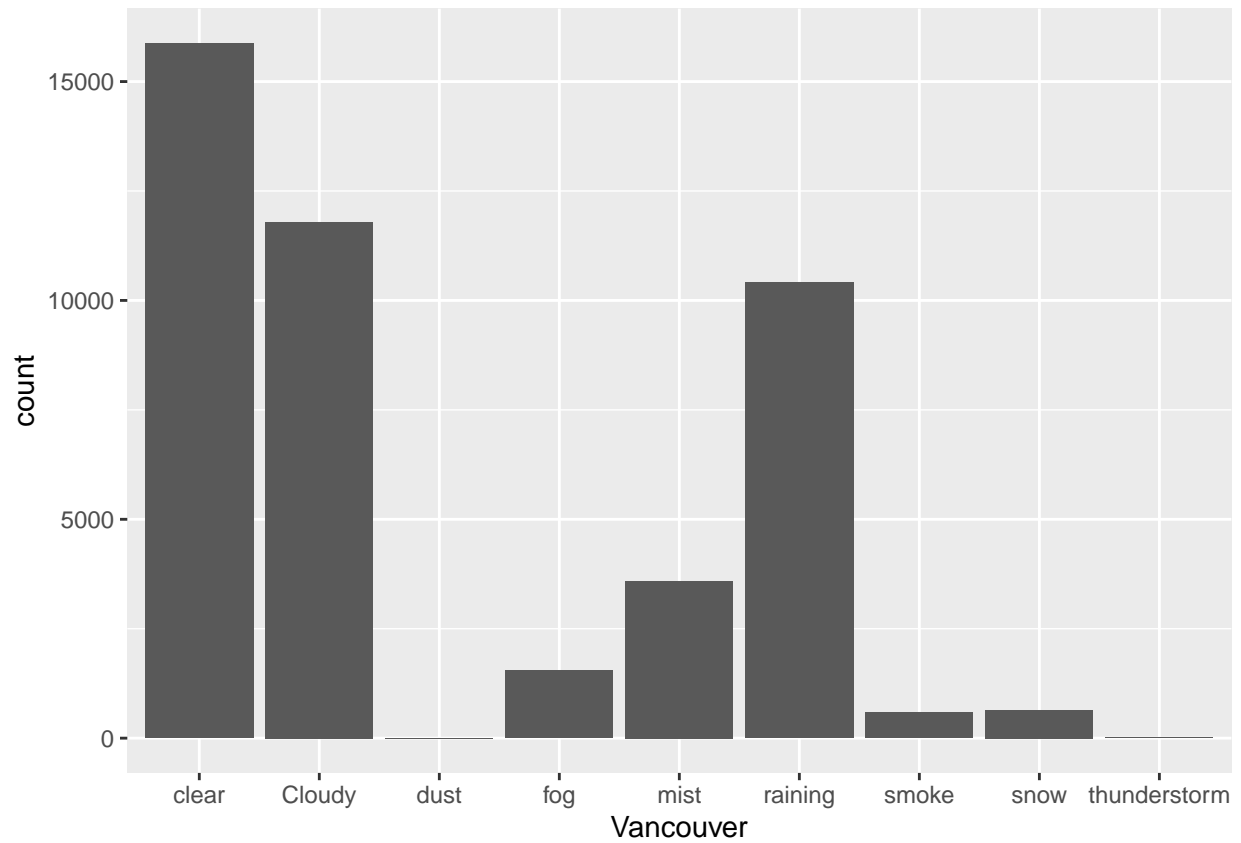
```
ggplot(airquality, aes(x=Temp)) + geom_histogram()
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```




## Bar plots

```
df_desc <- read.csv("../data/historical-hourly-weather-data/weather_description.csv")
ggplot(df_desc, aes(x=Vancouver)) + geom_bar()
```



```
ggplot(RetailSales, aes(x=Month, y=Sales)) + geom_bar(stat="identity")
```


```
## Warning: Removed 15 rows containing missing values ('position_stack()').
```



## Box plots


```
df_hum <- read.csv("../data/historical-hourly-weather-data/humidity.csv")
df_hum$datetime <- as.character(df_hum$datetime)
df_hum$month <- substr(df_hum$datetime, 6, 7)
ggplot(df_hum, aes(x=month, y=Vancouver)) + geom_boxplot()
```

```
## Warning: Removed 1826 rows containing non-finite values ('stat_boxplot()').
```




## Scatter plots and line plots

```
a = 3.4
v0 = 27
time <- runif(50, min=0, max=200)
distance <- sapply(time, function(x) v0 * x + 0.5 * a * x^2)
df <- data.frame(time,distance)
ggplot(df, aes(x=time, y=distance)) + geom_point()
```



```
ggplot(df, aes(x=time, y=distance)) + geom_line()
```




## Changing histogram defaults and adding aesthetics

```
df_hum <- read.csv("../data/historical-hourly-weather-data/humidity.csv")
ggplot(df_hum, aes(x=Vancouver)) + geom_histogram()
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```


```
## Warning: Removed 1826 rows containing non-finite values ('stat_bin()').
```






```
ggplot(df_hum, aes(x=Vancouver)) + geom_histogram(bins=15)
```

```
## Warning: Removed 1826 rows containing non-finite values ('stat_bin()').
```



```
ggplot(df_hum, aes(x=Vancouver)) + geom_histogram(bins=15, fill="white", color=1)
```


```
## Warning: Removed 1826 rows containing non-finite values ('stat_bin()').
```



```
ggplot(df_hum, aes(x=Vancouver)) +  
  geom_histogram(bins=15, fill="white", color=1) +  
  ggtitle("Humidity for Vancouver city") +  
  xlab("Humidity") +  
  theme(axis.text.x=element_text(size=12), axis.text.y=element_text(size=12))
```

```
## Warning: Removed 1826 rows containing non-finite values ('stat_bin()').
```


Humidity for Vancouver city



## Changing boxplot defaults and adding aesthetics

```
df_hum <- read.csv("../data/historical-hourly-weather-data/humidity.csv")
df_hum$datetime <- as.character(df_hum$datetime)
df_hum$month <- substr(df_hum$datetime, 6, 7)
ggplot(df_hum, aes(x=month, y=Vancouver)) +
  geom_boxplot(color=1, fill=3) +
  ylab("Humidity") +
  theme(axis.text.x=element_text(size=15),
        axis.text.y=element_text(size=15),
        axis.title.x=element_text(size=15, color=2),
        axis.title.y=element_text(size=15, color=2))
```

```
## Warning: Removed 1826 rows containing non-finite values ('stat_boxplot()').
```




## Part 2 - Grammar of graphics and visual components

### Layers

```
df <- read.csv("../data/gapminder-data.csv")
p1 <- ggplot(df, aes(x=Electricity_consumption_per_capita))
p2 <- p1 + geom_histogram()
p3 <- p1 + geom_histogram(bins=15)
p4 <- p3 + xlab("Electricity consumption per capita")
p2
```


```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
## Warning: Removed 1181 rows containing non-finite values ('stat_bin()').
```




p3

```
## Warning: Removed 1181 rows containing non-finite values ('stat_bin()').
```



p4

```
## Warning: Removed 1181 rows containing non-finite values ('stat_bin()').
```



## Scales


```
p1 <- ggplot(df, aes(x=gdp_per_capita))
p2 <- p1 + geom_histogram()
p3 <- p2 + scale_x_continuous(name='GDP per capita ($)',
                              limits=c(0, 50000),
                              breaks=seq(0, 40000, 4000),
                              labels=c('0', '4K', '8K', '12K', '16K', '20K',
                                        '24K', '28K', '32K', '36K', '40K'))
p3
```

```
## 'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.
```

```
## Warning: Removed 7 rows containing non-finite values ('stat_bin()').
```


```
## Warning: Removed 2 rows containing missing values ('geom_bar()').
```





## Polar coordinates


```
t <- seq(0, 360, by=15)
r <- 2
qplot(r, t) +
  coord_polar(theta="y") +
  scale_y_continuous(breaks=seq(0, 360, 30))
```



## Facets


```
p <- ggplot(df, aes(x=gdp_per_capita, y=Electricity_consumption_per_capita)) + geom_point()
p + facet_grid(Country ~ .)
```

```
## Warning: Removed 1181 rows containing missing values ('geom_point()').
```




```
p + facet_grid(. ~ Country)
```

```
## Warning: Removed 1181 rows containing missing values ('geom_point()').
```



```
p + facet_wrap(~Country)
```

```
## Warning: Removed 1181 rows containing missing values ('geom_point()').
```



## Shapes and colors

```
dfs <- subset(df, Country %in% c("Germany", "India", "China", "United States"))
var1 <- "Electricity_consumption_per_capita"
var2 <- "gdp_per_capita"
name1 <- "Electricity/capita"
name2 <- "GDP/capita"
p1 <- ggplot(dfs, aes_string(x=var1, y=var2)) +
  geom_point(color=2, shape=2) +
  xlim(0, 10000) + xlab(name1) + ylab(name2)
```

```
## Warning: 'aes_string()' was deprecated in ggplot2 3.0.0.
```

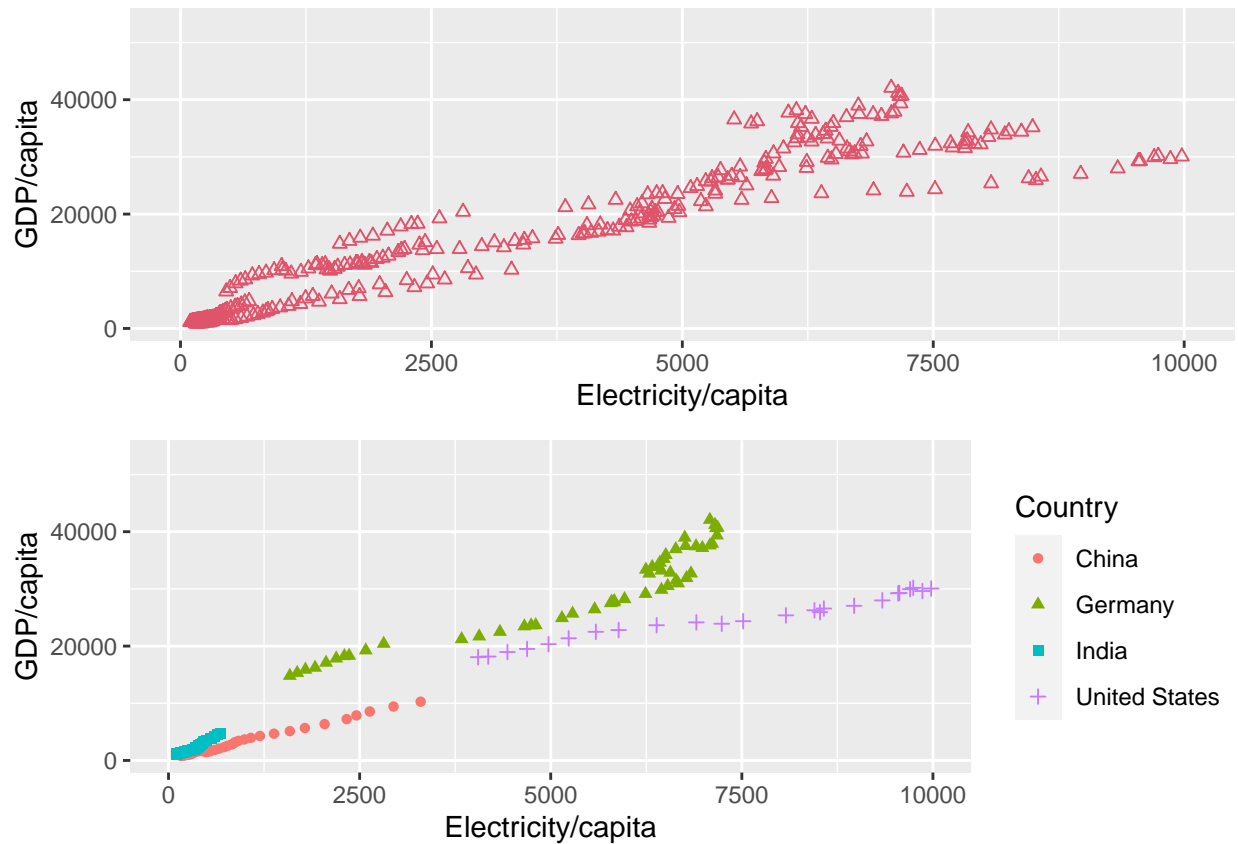
```
## Warning: Please use tidy evaluation idioms with 'aes()'.
```

```
## Warning: See also 'vignette("ggplot2-in-packages")' for more information.
```

```
p2 <- ggplot(dfs, aes_string(x=var1, y=var2)) +
  geom_point(aes(color=Country, shape=Country)) +
  xlim(0, 10000) + xlab(name1) + ylab(name2)
grid.arrange(p1, p2, nrow = 2)
```


```
## Warning: Removed 1209 rows containing missing values ('geom_point()').
```

```
## Warning: Removed 706 rows containing missing values ('geom_point()').
```




## Themes

```
dfn <- subset(HollywoodMovies,
  Genre %in% c("Action", "Adventure", "Comedy", "Drama",
    "Romantic Comedy") &
  LeadStudio %in% c("Lionsgate ", "Paramount Pictures ",
    "Twentieth Century Fox ", "Universal Pictures ", "Warner Bros. "))
p1 <- ggplot(dfn, aes(x=Genre, y=WorldGross))
p2 <- p1 + geom_bar(aes(fill=LeadStudio), stat="Identity", position="dodge")
p3 <- p2 + theme(axis.title.x=element_text(size=15),
  axis.title.y=element_text(size=15),
  plot.background=element_rect(fill="gray87"),
  panel.background=element_rect(fill="beige"),
  panel.grid.major=element_line(color="Gray", linetype=1))
p3
```

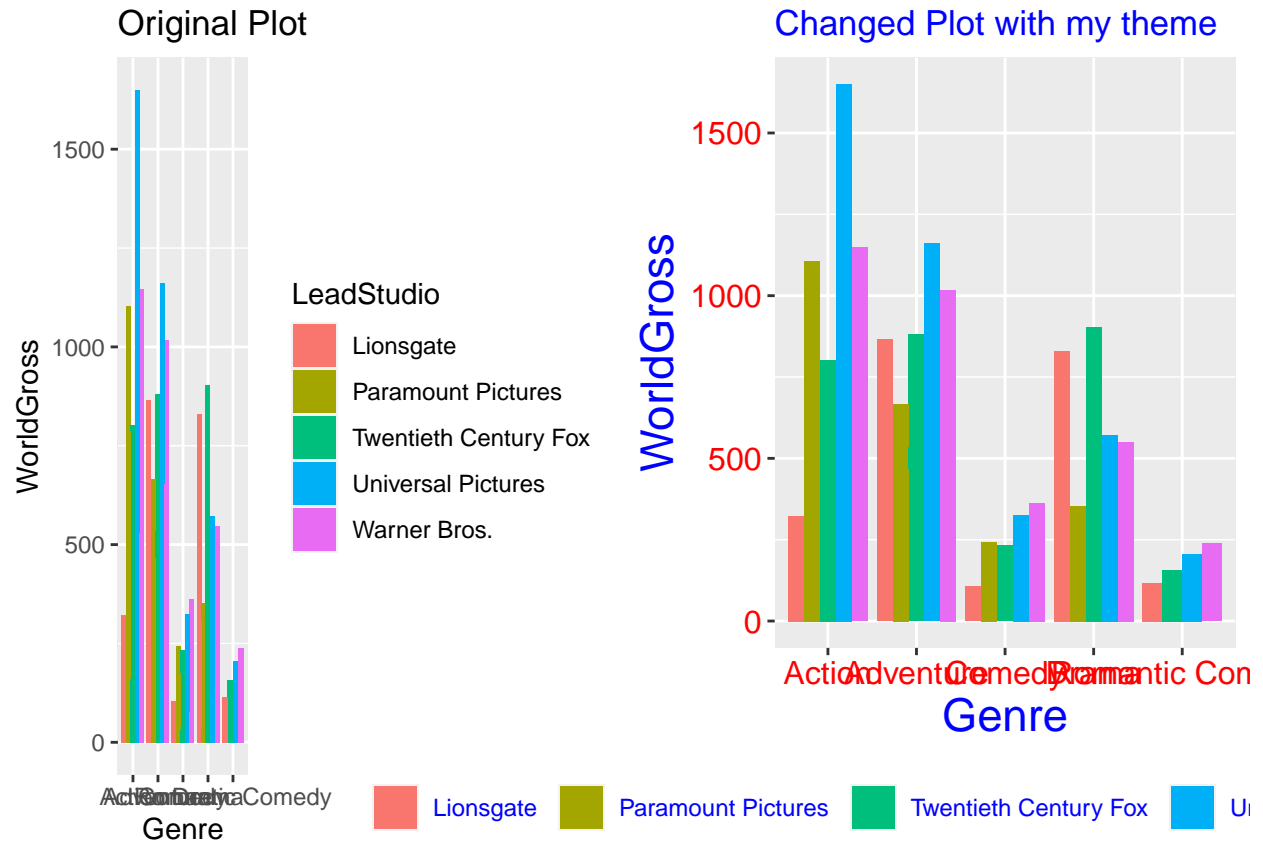


```
p4 <- p2 + theme_bw() + ggtitle("theme_bw()")
p5 <- p2 + theme_classic() + ggtitle("theme_classic()")
p6 <- p2 + theme_classic() + ggtitle("theme_gray()")
p7 <- p2 + theme_minimal() + ggtitle("theme_minimal()")
grid.arrange(p4, p5, p6, p7, nrow=2, ncol=2)
```




```
mytheme <- theme(legend.title=element_blank(),
  legend.position="bottom",
  text = element_text(color="Blue"),
  axis.text=element_text(size=12, color="Red"),
  axis.title=element_text(size=rel(1.5)))
p2 <- p2 + ggtitle("Original Plot")
p8 <- p2 + mytheme + ggtitle("Changed Plot with my theme")
grid.arrange(p2, p8, ncol=2)
```







```
p4 + scale_fill_brewer(palette="Spectral")
```



```
p4 + scale_fill_brewer(palette="Pastel1")
```



```
p4 + scale_fill_brewer(palette="Oranges")
```



## Part 3 - Advanced Geoms and Statistics

### Bubble charts


```
df <- read.csv("../data/gapminder-data.csv")
dfs <- subset(df, Country %in% c("Germany", "India", "China", "United States", "Japan"))

ggplot(dfs, aes(x=Year, y=Electricity_consumption_per_capita)) + geom_point(aes(size=population, color=
coord_cartesian(xlim=c(1950, 2020)) +
labs(subtitle="Electricity consumption vs Year", title="Bubble chart") +
ylab("Electricity consumption") +
scale_size(breaks=c(0, 1e+8, 0.3e+9, 0.5e+9, 1e+9, 1.5e+9), range=c(1, 5))
```

```
## Warning: Removed 842 rows containing missing values (‘geom_point()’).
```

## Bubble chart

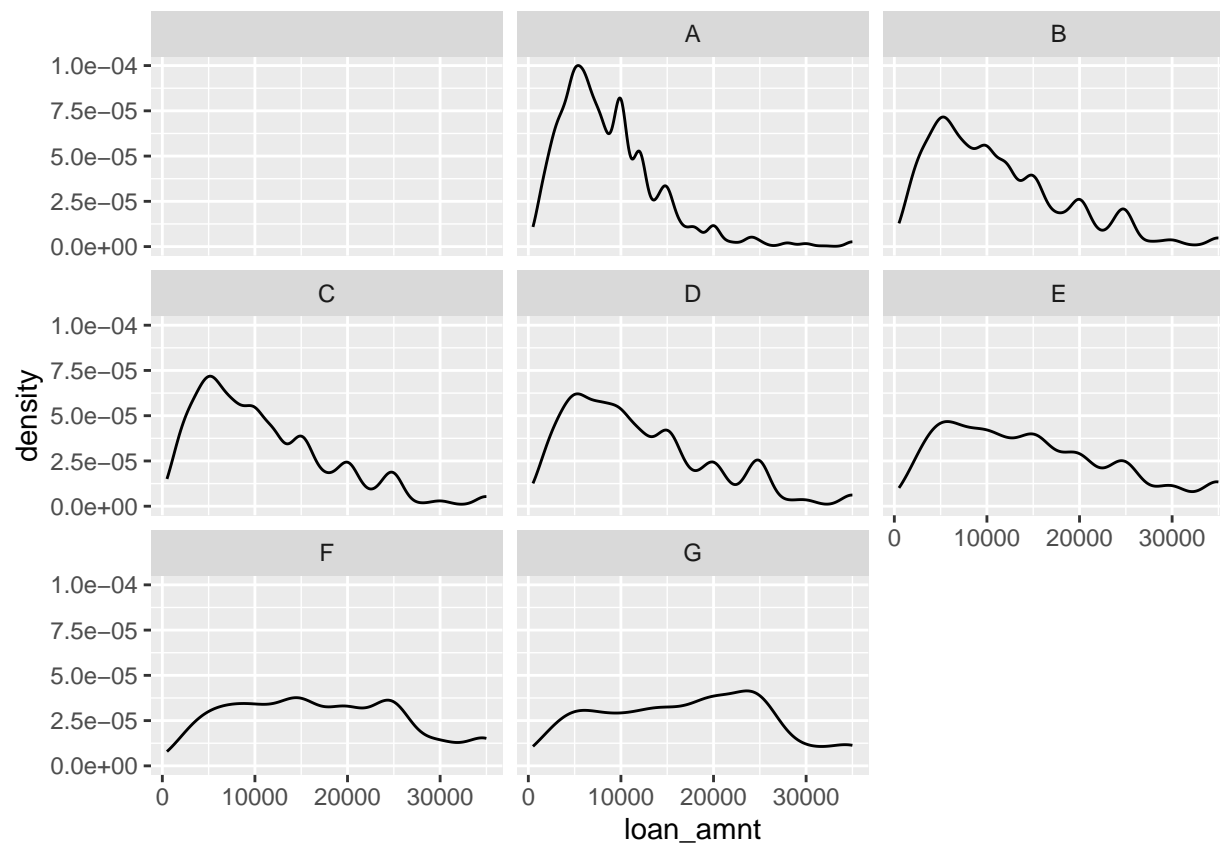
Electricity consumption vs Year



## Density plots


```
df <- read.csv("../data/LoanStats.csv")
ggplot(df, aes(x=loan_amnt)) + geom_density() + facet_wrap(~grade)
```

```
## Warning: Removed 7 rows containing non-finite values ('stat_density()').
```



```
df <- read.csv("../data/LoanStats.csv")
ggplot(df, aes(x=loan_amnt)) + geom_density(aes(fill=grade), alpha=1/2) +
scale_fill_brewer(palette="Dark2") + xlab("Loan Amount") + theme_light()
```

```
## Warning: Removed 7 rows containing non-finite values ('stat_density()').
```



## Time series

```
df_fb <- read.csv("../data/FB.csv")
df_fb$Date <- as.Date(df_fb$Date)
ggplot(df_fb, aes(x=Date, y=Close, group=1)) +
  geom_line(color="black", na.rm=TRUE) +
  ggtitle("Daily Closing Stock Prices: Facebook") +
  theme(plot.title = element_text(lineheight=.7, face="bold")) +
  scale_x_date(date_breaks='3 month') +
  theme_bw()
```

