

Tidyverse tutorial 1 - Basic operations

Ariane Ducellier

10/03/2023

Load R packages.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.3      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.3      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(skimr)
```

1. Read data (csv format)

Read csv files with basic R.

```
df <- read.csv("../data/gapminder-data.csv")
print(class(df))
```

```
## [1] "data.frame"
```

Read csv files with tidyverse.

```
df_t <- read_csv("../data/gapminder-data.csv")
```

```
## New names:
## Rows: 1512 Columns: 10
## -- Column specification
## ----- Delimiter: "," chr
## (1): Country dbl (9): ...1, Year, gdp_per_capita,
## Electricity_consumption_per_capita, und...
## i Use `spec()` to retrieve the full column specification for this data. i
## Specify the column types or set `show_col_types = FALSE` to quiet this message.
## * `` -> `...1`
```

```
print(class(df_t))
```

```
## [1] "spec_tbl_df" "tbl_df"      "tbl"         "data.frame"
```

```
df_t_sub <- read_csv("../data/gapminder-data.csv",
  col_select=c("Country", "Year", "gdp_per_capita"),
  na=c("", "NA"))

## New names:
## Rows: 1512 Columns: 3
## -- Column specification
## ----- Delimiter: "," chr
## (1): Country dbl (2): Year, gdp_per_capita
## i Use `spec()` to retrieve the full column specification for this data. i
## Specify the column types or set `show_col_types = FALSE` to quiet this message.
## * `` -> `...1`
```

2. Get basic information on the data.

Reads the first part of a data frame.

```
df_t <- read_csv("../data/gapminder-data.csv")

## New names:
## Rows: 1512 Columns: 10
## -- Column specification
## ----- Delimiter: "," chr
## (1): Country dbl (9): ...1, Year, gdp_per_capita,
## Electricity_consumption_per_capita, und...
## i Use `spec()` to retrieve the full column specification for this data. i
## Specify the column types or set `show_col_types = FALSE` to quiet this message.
## * `` -> `...1`

head(df_t, 3)
```

```
## # A tibble: 3 x 10
##   ...1 Country Year gdp_per_capita Electricity_consumption_p~1 under5mortality
##   <dbl> <chr>   <dbl>         <dbl>         <dbl>         <dbl>
## 1     0 Brazil  1800           1109             NA           417.
## 2     1 Brazil  1801           1109             NA           417.
## 3     2 Brazil  1802           1109             NA           417.
## # i abbreviated name: 1: Electricity_consumption_per_capita
## # i 4 more variables: Poverty <dbl>, BMI_male <dbl>, BMI_female <dbl>,
## #   population <dbl>
```

Reads the last part of a data frame.

```
tail(df_t, 3)

## # A tibble: 3 x 10
##   ...1 Country Year gdp_per_capita Electricity_consumpt~1 under5mortality
##   <dbl> <chr>   <dbl>         <dbl>         <dbl>         <dbl>
## 1 1509 United Stat~ 2013           51282             NA           6.9
## 2 1510 United Stat~ 2014           52118             NA           6.7
## 3 1511 United Stat~ 2015           53354             NA           6.5
## # i abbreviated name: 1: Electricity_consumption_per_capita
## # i 4 more variables: Poverty <dbl>, BMI_male <dbl>, BMI_female <dbl>,
## #   population <dbl>
```

Gets column specifications of a tibble.

```
spec(df_t)
```

```
## cols(  
##   ...1 = col_double(),  
##   Country = col_character(),  
##   Year = col_double(),  
##   gdp_per_capita = col_double(),  
##   Electricity_consumption_per_capita = col_double(),  
##   under5mortality = col_double(),  
##   Poverty = col_double(),  
##   BMI_male = col_double(),  
##   BMI_female = col_double(),  
##   population = col_double()  
## )
```

Prints the data: number of rows and columns, type of columns, and first rows.

```
glimpse(df_t)
```

```
## Rows: 1,512  
## Columns: 10  
## $ ...1          <dbl> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 1~  
## $ Country       <chr> "Brazil", "Brazil", "Brazil", "Braz~  
## $ Year          <dbl> 1800, 1801, 1802, 1803, 1804, 1805,~  
## $ gdp_per_capita <dbl> 1109, 1109, 1109, 1109, 1109, 1110,~  
## $ Electricity_consumption_per_capita <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~  
## $ under5mortality <dbl> 417.44, 417.44, 417.44, 417.44, 417~  
## $ Poverty       <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~  
## $ BMI_male      <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~  
## $ BMI_female    <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~  
## $ population    <dbl> 3639636, NA, NA, NA, NA, NA, NA, NA,~
```

Returns descriptive statistics on each column of a data frame.

```
summary(df_t)
```

```
##           ...1      Country      Year      gdp_per_capita  
## Min.      : 0.0    Length:1512    Min.      :1800    Min.      : 529  
## 1st Qu.: 377.8    Class :character 1st Qu.:1854    1st Qu.: 1124  
## Median : 755.5    Mode  :character  Median :1908    Median : 2496  
## Mean      : 755.5                Mean      :1908    Mean      : 7234  
## 3rd Qu.:1133.2                3rd Qu.:1961    3rd Qu.: 8219  
## Max.      :1511.0                Max.      :2015    Max.      :53354  
##  
## Electricity_consumption_per_capita under5mortality      Poverty  
## Min.      : 97.78                Min.      : 2.70    Min.      : 0.000  
## 1st Qu.: 1062.24                1st Qu.: 77.59    1st Qu.: 0.920  
## Median : 4310.62                Median :306.66    Median : 9.385  
## Mean      : 4386.74                Mean      :260.02    Mean      :15.338  
## 3rd Qu.: 6495.64                3rd Qu.:417.44    3rd Qu.:15.960  
## Max.      :13704.58                Max.      :539.16    Max.      :84.270  
## NA's      :1181                NA's      :1440  
## BMI_male      BMI_female      population  
## Min.      :20.62    Min.      :20.48    Min.      :3.640e+06  
## 1st Qu.:22.22    1st Qu.:21.90    1st Qu.:6.740e+07  
## Median :24.04    Median :24.57    Median :1.250e+08
```

```
## Mean      :24.16    Mean      :23.91    Mean      :2.996e+08
## 3rd Qu.   :26.17    3rd Qu.   :25.56    3rd Qu.   :3.767e+08
## Max.      :28.46    Max.      :28.34    Max.      :1.376e+09
## NA's      :1309     NA's      :1309     NA's      :945
```

Provides a broad overview of a data frame, handles data of all types, dispatching a different set of summary functions based on the types of columns in the data frame.

```
#skim(df_t)
```

3. The pipe operator

```
data(iris)
df_iris <- iris %>%
  group_by(Species) %>%
  summarize_if(is.numeric, mean) %>%
  ungroup() %>%
  gather(measure, value, -Species) %>%
  arrange(value)
```

```
data(iris)
df_iris_alt <- group_by(iris, Species)
df_iris_alt <- summarize_if(df_iris_alt, is.numeric, mean)
df_iris_alt <- ungroup(df_iris_alt)
df_iris_alt <- gather(df_iris_alt, measure, value, -Species)
df_iris_alt <- arrange(df_iris_alt, value)
```

4. Transform the data

```
header <- c("age", "workclass", "fnlwgt", "education",
  "education_num", "marital_status", "occupation",
  "relationship", "race", "sex", "capital_gain",
  "capital_loss", "hours_per_week", "native_country", "target")
df <- read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data",
  col_names=header, trim_ws=TRUE)
```

```
## Rows: 32561 Columns: 15
## -- Column specification -----
## Delimiter: ","
## chr (9): workclass, education, marital_status, occupation, relationship, rac...
## dbl (6): age, fnlwgt, education_num, capital_gain, capital_loss, hours_per_week
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

4.1 Slicing

Cuts unwanted parts of the data set.

```
df %>% slice_head(n=5)
```

```
## # A tibble: 5 x 15
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>         <dbl> <chr>         <chr>
```

```
## 1 39 State-gov 77516 Bachelors 13 Never-married Adm-cleri~
## 2 50 Self-emp-not-i~ 83311 Bachelors 13 Married-civ-s~ Exec-mana~
## 3 38 Private 215646 HS-grad 9 Divorced Handlers~~
## 4 53 Private 234721 11th 7 Married-civ-s~ Handlers~~
## 5 28 Private 338409 Bachelors 13 Married-civ-s~ Prof-spec~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

```
df %>% slice_tail(n=5)
```

```
## # A tibble: 5 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr> <dbl> <chr> <dbl> <chr> <chr>
## 1 27 Private 257302 Assoc-acdm 12 Married-civ-spo~ Tech-sup~
## 2 40 Private 154374 HS-grad 9 Married-civ-spo~ Machine-o~
## 3 58 Private 151910 HS-grad 9 Widowed Adm-cleri~
## 4 22 Private 201490 HS-grad 9 Never-married Adm-cleri~
## 5 52 Self-emp-inc 287927 HS-grad 9 Married-civ-spo~ Exec-mana~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

```
df %>% slice_min(age, prop=0.10)
```

```
## # A tibble: 3,895 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr> <dbl> <chr> <dbl> <chr> <chr>
## 1 17 ? 304873 10th 6 Never-married ?
## 2 17 Private 65368 11th 7 Never-married Sales
## 3 17 Private 245918 11th 7 Never-married Other-service
## 4 17 Private 191260 9th 5 Never-married Other-service
## 5 17 Private 270942 5th-6th 3 Never-married Other-service
## 6 17 Private 89821 11th 7 Never-married Other-service
## 7 17 Private 175024 11th 7 Never-married Handlers-clean~
## 8 17 ? 202521 11th 7 Never-married ?
## 9 17 ? 258872 11th 7 Never-married ?
## 10 17 Private 211870 9th 5 Never-married Other-service
## # i 3,885 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

```
df %>% slice_max(age, prop=0.30)
```

```
## # A tibble: 10,361 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr> <dbl> <chr> <dbl> <chr> <chr>
## 1 90 Private 51744 HS-grad 9 Never-married Other-ser~
## 2 90 Private 137018 HS-grad 9 Never-married Other-ser~
## 3 90 Private 221832 Bachelors 13 Married-civ-spo~ Exec-mana~
## 4 90 Private 52386 Some-college 10 Never-married Other-ser~
## 5 90 Private 171956 Some-college 10 Separated Adm-cleri~
## 6 90 Private 313986 11th 7 Never-married Handlers~~
## 7 90 ? 256514 Bachelors 13 Widowed ?
## 8 90 Private 52386 Some-college 10 Never-married Other-ser~
```

```
## 9 90 Private 141758 9th 5 Never-married Adm-cleri~
## 10 90 Local-gov 227796 Masters 14 Married-civ-spo~ Exec-mana~
## # i 10,351 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

```
df %>% slice_sample(n=10, replace=TRUE)
```

```
## # A tibble: 10 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>    <dbl> <chr>          <dbl> <chr>      <chr>
## 1 52 Private 218550 Some-college 10 Married-spouse~ Adm-cleri~
## 2 26 Private 173593 Masters 14 Married-civ-spo~ Machine-o~
## 3 29 State-gov 356089 Bachelors 13 Married-civ-spo~ Exec-mana~
## 4 34 Private 183473 HS-grad 9 Divorced Transport~
## 5 70 Private 145419 HS-grad 9 Widowed Adm-cleri~
## 6 22 Private 102632 HS-grad 9 Never-married Craft-rep~
## 7 39 Private 165215 Masters 14 Married-civ-spo~ Prof-spec~
## 8 20 Private 387779 11th 7 Never-married Transport~
## 9 28 Private 250135 Some-college 10 Divorced Exec-mana~
## 10 25 Private 281627 Masters 14 Never-married Prof-spec~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

4.2 Filtering

Apply a condition to one of the variables to filter unwanted rows of the data.

```
df %>% filter(age > 30)
```

```
## # A tibble: 21,989 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>    <dbl> <chr>          <dbl> <chr>      <chr>
## 1 39 State-gov 77516 Bachelors 13 Never-married Adm-cleri~
## 2 50 Self-emp-not~ 83311 Bachelors 13 Married-civ-s~ Exec-mana~
## 3 38 Private 215646 HS-grad 9 Divorced Handlers~
## 4 53 Private 234721 11th 7 Married-civ-s~ Handlers~
## 5 37 Private 284582 Masters 14 Married-civ-s~ Exec-mana~
## 6 49 Private 160187 9th 5 Married-spous~ Other-ser~
## 7 52 Self-emp-not~ 209642 HS-grad 9 Married-civ-s~ Exec-mana~
## 8 31 Private 45781 Masters 14 Never-married Prof-spec~
## 9 42 Private 159449 Bachelors 13 Married-civ-s~ Exec-mana~
## 10 37 Private 280464 Some-col~ 10 Married-civ-s~ Exec-mana~
## # i 21,979 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

4.3 Selecting

Select specific columns.

```
df %>% select(marital_status, age)
```

```
## # A tibble: 32,561 x 2
##   marital_status      age
##   <chr>             <dbl>
## 1 Never-married      39
## 2 Married-civ-spouse 50
## 3 Divorced           38
## 4 Married-civ-spouse 53
## 5 Married-civ-spouse 28
## 6 Married-civ-spouse 37
## 7 Married-spouse-absent 49
## 8 Married-civ-spouse 52
## 9 Never-married      31
## 10 Married-civ-spouse 42
## # i 32,551 more rows
```

4.4 Unique values

Get unique entries for categorical variables.

```
df %>% distinct(sex)
```

```
## # A tibble: 2 x 1
##   sex
##   <chr>
## 1 Male
## 2 Female
```

4.5 Grouping

Group by column and summarize.

```
df %>%
  group_by(workclass) %>%
  summarize(age_avg=mean(age))
```

```
## # A tibble: 9 x 2
##   workclass      age_avg
##   <chr>         <dbl>
## 1 ?            41.0
## 2 Federal-gov  42.6
## 3 Local-gov    41.8
## 4 Never-worked 20.6
## 5 Private      36.8
## 6 Self-emp-inc 46.0
## 7 Self-emp-not-inc 45.0
## 8 State-gov    39.4
## 9 Without-pay  47.8
```

4.6 Summarizing

The summary may be: - counting observations - counting available observations (i.e. not NA) - getting first or last value - compute statistics on each group (mean, standard deviation, quantile)

```
df %>% group_by(workclass) %>% summarize(n())
```

```
## # A tibble: 9 x 2
##   workclass      `n()`
```

```
##   <chr>           <int>
## 1 ?               1836
## 2 Federal-gov     960
## 3 Local-gov       2093
## 4 Never-worked    7
## 5 Private         22696
## 6 Self-emp-inc    1116
## 7 Self-emp-not-inc 2541
## 8 State-gov       1298
## 9 Without-pay     14
```

```
df %>% summarize(sum(!is.na(workclass)))
```

```
## # A tibble: 1 x 1
##   `sum(!is.na(workclass))`
##               <int>
## 1                 32561
```

```
df %>% group_by(workclass) %>% summarize(first(age))
```

```
## # A tibble: 9 x 2
##   workclass   `first(age)`
##   <chr>       <dbl>
## 1 ?          54
## 2 Federal-gov 35
## 3 Local-gov   56
## 4 Never-worked 18
## 5 Private     38
## 6 Self-emp-inc 47
## 7 Self-emp-not-inc 50
## 8 State-gov   39
## 9 Without-pay 65
```

```
df %>% group_by(workclass) %>% summarize(sd(capital_gain))
```

```
## # A tibble: 9 x 2
##   workclass   `sd(capital_gain)`
##   <chr>       <dbl>
## 1 ?          5147.
## 2 Federal-gov 4102.
## 3 Local-gov   5775.
## 4 Never-worked 0
## 5 Private     6424.
## 6 Self-emp-inc 17977.
## 7 Self-emp-not-inc 10986.
## 8 State-gov   3778.
## 9 Without-pay 1301.
```

```
df %>% group_by(workclass) %>% summarize(quantile(age, 0.5))
```

```
## # A tibble: 9 x 2
##   workclass   `quantile(age, 0.5)`
##   <chr>       <dbl>
## 1 ?          35
## 2 Federal-gov 43
## 3 Local-gov   41
## 4 Never-worked 18
```



```
## 5 Private 35
## 6 Self-emp-inc 45
## 7 Self-emp-not-inc 44
## 8 State-gov 39
## 9 Without-pay 57
```

We can also apply the summary over selected columns.

```
df %>% select(1, 3, 5, 11, 12, 13) %>% summarize(across(everything(), mean))

## # A tibble: 1 x 6
##   age   fnlwgt education_num capital_gain capital_loss hours_per_week
##   <dbl>   <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1  38.6 189778.           10.1          1078.           87.3          40.4
```

4.7 Arranging

To sort the data set.

```
df %>% arrange(native_country)

## # A tibble: 32,561 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>      <dbl> <chr>          <dbl> <chr>          <chr>
## 1  40 Private  121772 Assoc-voc      11 Married-civ-spo~ Craft-rep~
## 2  31 Private   84154 Some-college  10 Married-civ-spo~ Sales
## 3  18 Private  226956 HS-grad        9 Never-married   Other-ser~
## 4  32 ?       293936 7th-8th        4 Married-spouse~ ?
## 5  30 Private  117747 HS-grad        9 Married-civ-spo~ Sales
## 6  56 Private  203580 HS-grad        9 Married-civ-spo~ Adm-cleri~
## 7  45 Private  153141 HS-grad        9 Married-civ-spo~ Adm-cleri~
## 8  39 ?       157443 Masters       14 Married-civ-spo~ ?
## 9  34 State-gov 98101 Bachelors    13 Married-civ-spo~ Exec-mana~
## 10 42 Private  197583 Assoc-acdm    12 Married-civ-spo~ Exec-mana~
## # i 32,551 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

It is most useful to sort the data set after grouping and summarizing.

```
df %>%
  group_by(education) %>%
  summarize(count=n(),
            avg_net_gain = mean(capital_gain - capital_loss)) %>%
  arrange(desc(avg_net_gain))
```

```
## # A tibble: 16 x 3
##   education count avg_net_gain
##   <chr>      <int>      <dbl>
## 1 Prof-school  576      10183.
## 2 Doctorate   413      4507.
## 3 Masters    1723      2396.
## 4 Bachelors  5355      1638.
## 5 Preschool   51       832.
## 6 Assoc-voc  1382      642.
## 7 Assoc-acdm  1067      547.
```

```
## 8 Some-college 7291 527.
## 9 HS-grad 10501 506.
## 10 10th 933 348.
## 11 9th 514 313.
## 12 12th 433 252.
## 13 7th-8th 646 168.
## 14 11th 1175 165.
## 15 5th-6th 333 108.
## 16 1st-4th 168 77.5
```

4.8 Separating and uniting

This is often useful to create new columns.

```
df %>% separate(target, into=c("sign", "amount"), sep="\\b")
```

```
## Warning: Expected 2 pieces. Additional pieces discarded in 32561 rows [1, 2, 3, 4, 5, 6,
## 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, ...].
```

```
## # A tibble: 32,561 x 16
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>      <dbl> <chr>      <dbl> <chr>      <chr>
## 1 39 State-gov      77516 Bachelors      13 Never-married Adm-cleri~
## 2 50 Self-emp-not~ 83311 Bachelors      13 Married-civ-s~ Exec-mana~
## 3 38 Private      215646 HS-grad         9 Divorced      Handlers~
## 4 53 Private      234721 11th           7 Married-civ-s~ Handlers~
## 5 28 Private      338409 Bachelors      13 Married-civ-s~ Prof-spec~
## 6 37 Private      284582 Masters       14 Married-civ-s~ Exec-mana~
## 7 49 Private      160187 9th            5 Married-spous~ Other-ser~
## 8 52 Self-emp-not~ 209642 HS-grad         9 Married-civ-s~ Exec-mana~
## 9 31 Private      45781 Masters       14 Never-married Prof-spec~
## 10 42 Private      159449 Bachelors      13 Married-civ-s~ Exec-mana~
## # i 32,551 more rows
## # i 9 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, sign <chr>, amount <chr>
```

```
df %>% unite(sex, race, age, col="description", sep="_", remove=FALSE)
```

```
## # A tibble: 32,561 x 16
##   description      age workclass fnlwgt education education_num marital_status
##   <chr>      <dbl> <chr>      <dbl> <chr>      <dbl> <chr>
## 1 Male_White_39    39 State-gov  77516 Bachelors      13 Never-married
## 2 Male_White_50    50 Self-emp~  83311 Bachelors      13 Married-civ-s~
## 3 Male_White_38    38 Private  215646 HS-grad         9 Divorced
## 4 Male_Black_53    53 Private  234721 11th           7 Married-civ-s~
## 5 Female_Black_28   28 Private  338409 Bachelors      13 Married-civ-s~
## 6 Female_White_37   37 Private  284582 Masters       14 Married-civ-s~
## 7 Female_Black_49   49 Private  160187 9th            5 Married-spous~
## 8 Male_White_52    52 Self-emp~ 209642 HS-grad         9 Married-civ-s~
## 9 Female_White_31   31 Private   45781 Masters       14 Never-married
## 10 Male_White_42   42 Private  159449 Bachelors      13 Married-civ-s~
## # i 32,551 more rows
## # i 9 more variables: occupation <chr>, relationship <chr>, race <chr>,
## #   sex <chr>, capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

4.8 Mutate function

Designed to create new variables.

```
df %>%
  mutate(total_gain = capital_gain - capital_loss,
         tax = ifelse(total_gain >= 15000,
                      total_gain * 0.1,
                      0)
  ) %>%
  arrange(desc(tax))

## # A tibble: 32,561 x 17
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>          <dbl> <chr>         <chr>
## 1    54 Self-emp-inc  166459 Prof-sch~      15 Married-civ-s~ Prof-spec~
## 2    52 Private      152234 HS-grad         9 Married-civ-s~ Exec-mana~
## 3    53 Self-emp-inc  263925 HS-grad         9 Married-civ-s~ Sales
## 4    52 Private      118025 Bachelors      13 Married-civ-s~ Exec-mana~
## 5    46 Private      370119 Prof-sch~      15 Married-civ-s~ Prof-spec~
## 6    43 Private      176270 Bachelors      13 Married-civ-s~ Exec-mana~
## 7    49 Private      159816 Bachelors      13 Married-civ-s~ Prof-spec~
## 8    50 Private      171338 Some-col~     10 Married-civ-s~ Exec-mana~
## 9    22 Self-emp-not-- 202920 HS-grad         9 Never-married Prof-spec~
## 10   43 Self-emp-inc  172826 Some-col~     10 Married-civ-s~ Sales
## # i 32,551 more rows
## # i 10 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>, total_gain <dbl>, tax <dbl>
```

We can use the mutate function to replace values (e.g. “?” by NA).

```
for (variable in colnames(df)) {
  print (
    paste (variable, dim(drop_na(df[df[variable]=="?", variable]))[1])
  )
}

## [1] "age 0"
## [1] "workclass 1836"
## [1] "fnlwgt 0"
## [1] "education 0"
## [1] "education_num 0"
## [1] "marital_status 0"
## [1] "occupation 1843"
## [1] "relationship 0"
## [1] "race 0"
## [1] "sex 0"
## [1] "capital_gain 0"
## [1] "capital_loss 0"
## [1] "hours_per_week 0"
## [1] "native_country 583"
## [1] "target 0"
```

```
df_replaced <- df %>%
  mutate(workclass = replace(workclass, workclass=="?", NA),
         occupation = replace(occupation, occupation=="?", NA),
```

```

    native_country =replace(native_country, native_country=="?", NA)
  )

for (variable in colnames(df_replaced)) {
  print (
    paste (variable, dim(drop_na(df_replaced[df_replaced[variable]=="?", variable]))[1])
  )
}

```

```

## [1] "age 0"
## [1] "workclass 0"
## [1] "fnlwgt 0"
## [1] "education 0"
## [1] "education_num 0"
## [1] "marital_status 0"
## [1] "occupation 0"
## [1] "relationship 0"
## [1] "race 0"
## [1] "sex 0"
## [1] "capital_gain 0"
## [1] "capital_loss 0"
## [1] "hours_per_week 0"
## [1] "native_country 0"
## [1] "target 0"

```

```

df %>% mutate(workclass = na_if(workclass, "?"),
              occupation = na_if(occupation, "?"),
              native_country = na_if(native_country, "?"))

```

```

## # A tibble: 32,561 x 15
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>          <dbl> <chr>         <chr>
## 1    39 State-gov      77516 Bachelors         13 Never-married Adm-cleri-
## 2    50 Self-emp-not-- 83311 Bachelors         13 Married-civ-s~ Exec-mana-
## 3    38 Private       215646 HS-grad           9 Divorced      Handlers--
## 4    53 Private       234721 11th              7 Married-civ-s~ Handlers--
## 5    28 Private       338409 Bachelors         13 Married-civ-s~ Prof-spec~
## 6    37 Private       284582 Masters          14 Married-civ-s~ Exec-mana-
## 7    49 Private       160187 9th                5 Married-spous~ Other-ser~
## 8    52 Self-emp-not-- 209642 HS-grad           9 Married-civ-s~ Exec-mana-
## 9    31 Private       45781 Masters          14 Never-married Prof-spec~
## 10   42 Private       159449 Bachelors         13 Married-civ-s~ Exec-mana-
## # i 32,551 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>

```

Additional examples:

```

df %>%
  mutate(over_under = recode(target, "<=50K"="under",
                              ">50K"="over")) %>%
  select(target, over_under)

```

```

## # A tibble: 32,561 x 2
##   target over_under

```

```
##      <chr>  <chr>
## 1 <=50K  under
## 2 <=50K  under
## 3 <=50K  under
## 4 <=50K  under
## 5 <=50K  under
## 6 <=50K  under
## 7 <=50K  under
## 8 >50K   over
## 9 >50K   over
## 10 >50K  over
## # i 32,551 more rows

df %>%
  mutate(age_avg = mean(age),
         over_under_age_avg = cut(age,
                                c(0, mean(age), max(age)),
                                c("Lower than avg", "Above the avg")))
  ) %>%
  select(age, age_avg, over_under_age_avg)

## # A tibble: 32,561 x 3
##       age age_avg over_under_age_avg
##   <dbl>   <dbl> <fct>
## 1    39    38.6 Above the avg
## 2    50    38.6 Above the avg
## 3    38    38.6 Lower than avg
## 4    53    38.6 Above the avg
## 5    28    38.6 Lower than avg
## 6    37    38.6 Lower than avg
## 7    49    38.6 Above the avg
## 8    52    38.6 Above the avg
## 9    31    38.6 Lower than avg
## 10   42    38.6 Above the avg
## # i 32,551 more rows
```

4.9 Joining tibbles

```
sales <- data.frame(
  date = c("2022-01-01", "2022-01-02", "2022-01-03", "2022-01-04", "2022-01-05"),
  store_cd= c(1, 2, 3, 4, 5),
  product_cd= c(1, 2, 3, 4, 5),
  qty= c(10, 12, 9, 12, 8),
  sales= c(30, 60, 45, 24, 32)
)

stores <- data.frame(
  store_cd= c(1, 2, 3, 4, 6),
  address= c("1 main st", "20 side st", "19 square blvd", "101 first st", "1002 retail ave"),
  city= c("Main", "East", "West", "North", "South"),
  open_hours= c("7-23", "7-23", "9-21", "9-21", "9-21")
)

products <- data.frame(
```

```
product_cd= c(1, 2, 3, 4, 6),
description= c("Soft drink", "Frozen snack", "Fruit", "Water", "Fruit 2"),
unit_price= c(3.0, 5.0, 5.0, 2.0, 4.0),
unit_measure= c("each", "each", "kg", "each", "kg")
)
```

4.9.1 Left join

All the rows from sales and matched rows from products.

```
sales %>% left_join(products, by="product_cd")
```

```
##      date store_cd product_cd qty sales description unit_price unit_measure
## 1 2022-01-01      1          1  10    30   Soft drink          3         each
## 2 2022-01-02      2          2  12    60  Frozen snack          5         each
## 3 2022-01-03      3          3   9    45      Fruit           5          kg
## 4 2022-01-04      4          4  12    24      Water           2         each
## 5 2022-01-05      5          5   8    32      <NA>          NA        <NA>
```

4.9.2 Right join

All the rows from stores and matched rows from sales.

```
sales %>% right_join(stores, by="store_cd")
```

```
##      date store_cd product_cd qty sales address city open_hours
## 1 2022-01-01      1          1  10    30   1 main st Main      7-23
## 2 2022-01-02      2          2  12    60   20 side st East      7-23
## 3 2022-01-03      3          3   9    45  19 square blvd West     9-21
## 4 2022-01-04      4          4  12    24  101 first st North    9-21
## 5      <NA>      6          NA   NA    NA 1002 retail ave South    9-21
```

4.9.3 Inner join

All the rows common to sales and stores.

```
sales %>% inner_join(stores, by="store_cd")
```

```
##      date store_cd product_cd qty sales address city open_hours
## 1 2022-01-01      1          1  10    30   1 main st Main      7-23
## 2 2022-01-02      2          2  12    60   20 side st East      7-23
## 3 2022-01-03      3          3   9    45  19 square blvd West     9-21
## 4 2022-01-04      4          4  12    24  101 first st North    9-21
```

4.9.4 Full join

All the rows from sales and stores.

```
sales %>% full_join(stores)
```

```
## Joining with `by = join_by(store_cd)`
```

```
##      date store_cd product_cd qty sales address city open_hours
## 1 2022-01-01      1          1  10    30   1 main st Main      7-23
## 2 2022-01-02      2          2  12    60   20 side st East      7-23
## 3 2022-01-03      3          3   9    45  19 square blvd West     9-21
## 4 2022-01-04      4          4  12    24  101 first st North    9-21
## 5 2022-01-05      5          5   8    32      <NA> <NA>      <NA>
```

```
## 6      <NA>      6      NA  NA      NA 1002 retail ave South      9-21
```

4.9.5 Anti-join

Only rows that are in sales but not in products.

```
sales %>% anti_join(products)
```

```
## Joining with `by = join_by(product_cd)`
```

```
##      date store_cd product_cd qty sales
## 1 2022-01-05      5          5    8    32
```

4.10 Reshaping tables

```
df_wide <- data.frame(
  project = c("project1", "project2", "project3"),
  Jan= sample(1000:2000, 3),
  Feb= sample(1000:2000, 3),
  Mar= sample(1000:2000, 3)
)
```

This is not a tidy data set (there is more than 1 observation per row).

```
df_long <- df_wide %>%
  pivot_longer(cols= 2:4,
               names_to = "months",
               values_to = "expenses")
```

If we need a smaller table for visualization or for a presentation.

```
df_wide_2 <- df_long %>%
  pivot_wider(names_from = "months",
              values_from = "expenses")
```