Tidyverse tutorial 1 - Basic operations

Ariane Ducellier

10/03/2023

Load R packages.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----------------------- tidyverse 2.0.0 --
## v dplyr     1.1.3     v readr     2.1.4
## v forcats   1.0.0     v stringr   1.5.0
## v ggplot2   3.4.3     v tibble    3.2.1
## v lubridate 1.9.2     v tidyr     1.3.0
## v purrr     1.0.2
## -- Conflicts ----------------------------------------- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(skimr)
```

1. Read data (csv format)

Read csv files with basic R.

```
df <- read.csv("../data/gapminder-data.csv")
print(class(df))
```

```
## [1] "data.frame"
```

Read csv files with tidyverse.

```
df_t <- read_csv("../data/gapminder-data.csv")
```

```
## New names:
## Rows: 1512 Columns: 10
## -- Column specification
## ----------------------------------------------------------- Delimiter: "," chr
## (1): Country dbl (9): ...1, Year, gdp_per_capita,
## Electricity_consumption_per_capita, und...
## i Use 'spec()' to retrieve the full column specification for this data. i
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.
## * '' -> '...1'
```

```
print(class(df_t))
```

```
## [1] "spec_tbl_df" "tbl_df"       "tbl"          "data.frame"
```

```
df_t_sub <- read_csv("../data/gapminder-data.csv",
  col_select=c("Country", "Year", "gdp_per_capita"),
  na=c("", "NA"))
```

```
## New names:
## Rows: 1512 Columns: 3
## -- Column specification
## ---------------------------------------------------------- Delimiter: "," chr
## (1): Country dbl (2): Year, gdp_per_capita
## i Use 'spec()' to retrieve the full column specification for this data. i
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.
## * '' -> '...1'
```

2. Get basic information on the data.

Reads the first part of a data frame.

```
df_t <- read_csv("../data/gapminder-data.csv")
```

```
## New names:
## Rows: 1512 Columns: 10
## -- Column specification
## ---------------------------------------------------------- Delimiter: "," chr
## (1): Country dbl (9): ...1, Year, gdp_per_capita,
## Electricity_consumption_per_capita, und...
## i Use 'spec()' to retrieve the full column specification for this data. i
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.
## * '' -> '...1'
```

```
head(df_t, 3)
```

```
## # A tibble: 3 x 10
##     ...1 Country  Year gdp_per_capita Electricity_consumption_p~1 under5mortality
##    <dbl> <chr>   <dbl>          <dbl>                       <dbl>           <dbl>
## 1      0 Brazil   1800           1109                          NA            417.
## 2      1 Brazil   1801           1109                          NA            417.
## 3      2 Brazil   1802           1109                          NA            417.
## # i abbreviated name: 1: Electricity_consumption_per_capita
## # i 4 more variables: Poverty <dbl>, BMI_male <dbl>, BMI_female <dbl>,
## #   population <dbl>
```

Reads the last part of a data frame.

```
tail(df_t, 3)
```

```
## # A tibble: 3 x 10
##    ...1 Country      Year gdp_per_capita Electricity_consumpt~1 under5mortality
##   <dbl> <chr>       <dbl>          <dbl>                  <dbl>           <dbl>
## 1  1509 United Stat~ 2013          51282                     NA             6.9
## 2  1510 United Stat~ 2014          52118                     NA             6.7
## 3  1511 United Stat~ 2015          53354                     NA             6.5
## # i abbreviated name: 1: Electricity_consumption_per_capita
## # i 4 more variables: Poverty <dbl>, BMI_male <dbl>, BMI_female <dbl>,
## #   population <dbl>
```

Gets column specifications of a tibble.

```
spec(df_t)
```

```
## cols(
##   ...1 = col_double(),
##   Country = col_character(),
##   Year = col_double(),
##   gdp_per_capita = col_double(),
##   Electricity_consumption_per_capita = col_double(),
##   under5mortality = col_double(),
##   Poverty = col_double(),
##   BMI_male = col_double(),
##   BMI_female = col_double(),
##   population = col_double()
## )
```

Prints the data: number of rows and columns, type of columns, and first rows.

```
glimpse(df_t)
```

```
## Rows: 1,512
## Columns: 10
## $ ...1                               <dbl> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 1~
## $ Country                            <chr> "Brazil", "Brazil", "Brazil", "Braz~
## $ Year                               <dbl> 1800, 1801, 1802, 1803, 1804, 1805,~
## $ gdp_per_capita                     <dbl> 1109, 1109, 1109, 1109, 1109, 1110,~
## $ Electricity_consumption_per_capita <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ under5mortality                    <dbl> 417.44, 417.44, 417.44, 417.44, 417~
## $ Poverty                            <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ BMI_male                           <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ BMI_female                         <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ population                         <dbl> 3639636, NA, NA, NA, NA, NA, NA, NA~
```

Returns descriptive statistics on each column of a data frame.

```
summary(df_t)
```

```
##        ...1              Country                Year       gdp_per_capita
##   Min.   :   0.0   Length:1512        Min.   :1800   Min.   :  529
##   1st Qu.: 377.8   Class :character   1st Qu.:1854   1st Qu.: 1124
##   Median : 755.5   Mode  :character   Median :1908   Median : 2496
##   Mean   : 755.5                      Mean   :1908   Mean   : 7234
##   3rd Qu.:1133.2                      3rd Qu.:1961   3rd Qu.: 8219
##   Max.   :1511.0                      Max.   :2015   Max.   :53354
##
##   Electricity_consumption_per_capita under5mortality     Poverty
##   Min.   :   97.78                   Min.   :  2.70   Min.   : 0.000
##   1st Qu.: 1062.24                   1st Qu.: 77.59   1st Qu.: 0.920
##   Median : 4310.62                   Median :306.66   Median : 9.385
##   Mean   : 4386.74                   Mean   :260.02   Mean   :15.338
##   3rd Qu.: 6495.64                   3rd Qu.:417.44   3rd Qu.:15.960
##   Max.   :13704.58                   Max.   :539.16   Max.   :84.270
##   NA's   :1181                                        NA's   :1440
##     BMI_male        BMI_female      population
##   Min.   :20.62   Min.   :20.48   Min.   :3.640e+06
##   1st Qu.:22.22   1st Qu.:21.90   1st Qu.:6.740e+07
##   Median :24.04   Median :24.57   Median :1.250e+08
##   Mean   :24.16   Mean   :23.91   Mean   :2.996e+08
##   3rd Qu.:26.17   3rd Qu.:25.56   3rd Qu.:3.767e+08
##   Max.   :28.46   Max.   :28.34   Max.   :1.376e+09
##   NA's   :1309    NA's   :1309    NA's   :945
```

Provides a broad overview of a data frame, handles data of all types, dispatching a different set of summary functions based on the types of columns in the data frame.

```
#skim(df_t)
```

3. The pipe operator

```
data(iris)
df_iris <- iris %>%
  group_by(Species) %>%
  summarize_if(is.numeric, mean) %>%
  ungroup() %>%
  gather(measure, value, -Species) %>%
  arrange(value)
```

```
data(iris)
df_iris_alt <- group_by(iris, Species)
df_iris_alt <- summarize_if(df_iris_alt, is.numeric, mean)
df_iris_alt <- ungroup(df_iris_alt)
df_iris_alt <- gather(df_iris_alt, measure, value, -Species)
df_iris_alt <- arrange(df_iris_alt, value)
```

4. Transform the data

```r
header <- c("age", "workclass", "fnlwgt", "education",
  "education_num", "marital_status", "occupation",
  "relationship", "race", "sex", "capital_gain",
  "capital_loss", "hours_per_week", "native_country", "target")
df <- read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data",
  col_names=header, trim_ws=TRUE)
```

```
## Rows: 32561 Columns: 15
## -- Column specification ---------------------------------------------------
## Delimiter: ","
## chr (9): workclass, education, marital_status, occupation, relationship, rac...
## dbl (6): age, fnlwgt, education_num, capital_gain, capital_loss, hours_per_week
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

4.1 Slicing

Cuts unwanted parts of the data set.

```r
df %>% slice_head(n=5)
```

```
## # A tibble: 5 x 15
##     age workclass       fnlwgt education education_num marital_status occupation
##   <dbl> <chr>            <dbl> <chr>             <dbl> <chr>          <chr>
## 1    39 State-gov        77516 Bachelors            13 Never-married  Adm-cleri~
## 2    50 Self-emp-not-i~  83311 Bachelors            13 Married-civ-s~ Exec-mana~
## 3    38 Private         215646 HS-grad               9 Divorced       Handlers-~
## 4    53 Private         234721 11th                  7 Married-civ-s~ Handlers-~
## 5    28 Private         338409 Bachelors            13 Married-civ-s~ Prof-spec~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```r
df %>% slice_tail(n=5)
```

```
## # A tibble: 5 x 15
##     age workclass      fnlwgt education education_num marital_status  occupation
##   <dbl> <chr>           <dbl> <chr>             <dbl> <chr>           <chr>
## 1    27 Private        257302 Assoc-acdm           12 Married-civ-spo~ Tech-supp~
## 2    40 Private        154374 HS-grad               9 Married-civ-spo~ Machine-o~
## 3    58 Private        151910 HS-grad               9 Widowed          Adm-cleri~
## 4    22 Private        201490 HS-grad               9 Never-married    Adm-cleri~
## 5    52 Self-emp-inc   287927 HS-grad               9 Married-civ-spo~ Exec-mana~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_min(age, prop=0.10)
```

```
## # A tibble: 3,895 x 15
##       age workclass  fnlwgt education education_num marital_status occupation
##     <dbl> <chr>       <dbl> <chr>             <dbl> <chr>          <chr>
## 1     17 ?          304873 10th                  6 Never-married  ?
## 2     17 Private     65368 11th                  7 Never-married  Sales
## 3     17 Private    245918 11th                  7 Never-married  Other-service
## 4     17 Private    191260 9th                   5 Never-married  Other-service
## 5     17 Private    270942 5th-6th               3 Never-married  Other-service
## 6     17 Private     89821 11th                  7 Never-married  Other-service
## 7     17 Private    175024 11th                  7 Never-married  Handlers-clean~
## 8     17 ?          202521 11th                  7 Never-married  ?
## 9     17 ?          258872 11th                  7 Never-married  ?
## 10    17 Private    211870 9th                   5 Never-married  Other-service
## # i 3,885 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_max(age, prop=0.30)
```

```
## # A tibble: 10,361 x 15
##       age workclass  fnlwgt education    education_num marital_status    occupation
##     <dbl> <chr>       <dbl> <chr>                <dbl> <chr>             <chr>
## 1     90 Private     51744 HS-grad                  9 Never-married     Other-ser~
## 2     90 Private    137018 HS-grad                  9 Never-married     Other-ser~
## 3     90 Private    221832 Bachelors               13 Married-civ-spo~  Exec-mana~
## 4     90 Private     52386 Some-college            10 Never-married     Other-ser~
## 5     90 Private    171956 Some-college            10 Separated         Adm-cleri~
## 6     90 Private    313986 11th                     7 Never-married     Handlers-~
## 7     90 ?          256514 Bachelors               13 Widowed           ?
## 8     90 Private     52386 Some-college            10 Never-married     Other-ser~
## 9     90 Private    141758 9th                      5 Never-married     Adm-cleri~
## 10    90 Local-gov  227796 Masters                 14 Married-civ-spo~  Exec-mana~
## # i 10,351 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_sample(n=10, replace=TRUE)
```

```
## # A tibble: 10 x 15
##       age workclass    fnlwgt education    education_num marital_status occupation
##     <dbl> <chr>         <dbl> <chr>                <dbl> <chr>          <chr>
## 1     50 Private      128378 5th-6th                  3 Married-civ-s~ Handlers-~
## 2     52 Federal-gov  207841 Some-college            10 Married-civ-s~ Adm-cleri~
## 3     31 Federal-gov  166626 HS-grad                  9 Married-civ-s~ Adm-cleri~
## 4     23 Private      256278 7th-8th                  4 Married-civ-s~ Handlers-~
## 5     23 ?            218415 Some-college            10 Never-married  ?
## 6     32 Private      216361 Some-college            10 Married-civ-s~ Protectiv~
```

```
## 7     25 Private         55718 Some-college         10 Never-married  Tech-supp~
## 8     26 Private        150132 HS-grad               9 Never-married  Adm-cleri~
## 9     47 Private        189680 Bachelors            13 Married-civ-s~ Sales
## 10    30 Private        229895 10th                  6 Divorced       Other-ser~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

4.2 Filtering

Apply a condition to one of the variables to filter unwanted rows of the data.

```
df %>% filter(age > 30)
```

```
## # A tibble: 21,989 x 15
##       age workclass        fnlwgt education education_num marital_status occupation
##     <dbl> <chr>             <dbl> <chr>            <dbl> <chr>          <chr>
## 1      39 State-gov         77516 Bachelors           13 Never-married  Adm-cleri~
## 2      50 Self-emp-not-~    83311 Bachelors           13 Married-civ-s~ Exec-mana~
## 3      38 Private          215646 HS-grad              9 Divorced       Handlers-~
## 4      53 Private          234721 11th                 7 Married-civ-s~ Handlers-~
## 5      37 Private          284582 Masters             14 Married-civ-s~ Exec-mana~
## 6      49 Private          160187 9th                  5 Married-spous~ Other-ser~
## 7      52 Self-emp-not-~   209642 HS-grad              9 Married-civ-s~ Exec-mana~
## 8      31 Private           45781 Masters             14 Never-married  Prof-spec~
## 9      42 Private          159449 Bachelors           13 Married-civ-s~ Exec-mana~
## 10     37 Private          280464 Some-col~           10 Married-civ-s~ Exec-mana~
## # i 21,979 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

4.3 Selecting

Select specific columns.

```
df %>% select(marital_status, age)
```

```
## # A tibble: 32,561 x 2
##    marital_status         age
##    <chr>                <dbl>
## 1 Never-married           39
## 2 Married-civ-spouse      50
## 3 Divorced                38
## 4 Married-civ-spouse      53
## 5 Married-civ-spouse      28
## 6 Married-civ-spouse      37
## 7 Married-spouse-absent   49
## 8 Married-civ-spouse      52
## 9 Never-married           31
## 10 Married-civ-spouse     42
## # i 32,551 more rows
```

4.4 Unique values

Get unique entries for categorical variables.

```
df %>% distinct(sex)
```

```
## # A tibble: 2 x 1
##   sex
##   <chr>
## 1 Male
## 2 Female
```

4.5 Grouping

Group by column and summarize.

```
df %>%
  group_by(workclass) %>%
  summarize(age_avg=mean(age))
```

```
## # A tibble: 9 x 2
##   workclass       age_avg
##   <chr>             <dbl>
## 1 ?                  41.0
## 2 Federal-gov        42.6
## 3 Local-gov          41.8
## 4 Never-worked       20.6
## 5 Private            36.8
## 6 Self-emp-inc       46.0
## 7 Self-emp-not-inc   45.0
## 8 State-gov          39.4
## 9 Without-pay        47.8
```

4.6 Summarizing

The summary may be: - counting observations - counting available observations (i.e. not NA) - getting first
or last value - compute statistics on each group (mean, standard deviation, quantile)

```
df %>% group_by(workclass) %>% summarize(n())
```

```
## # A tibble: 9 x 2
##   workclass       `n()`
##   <chr>           <int>
## 1 ?                1836
## 2 Federal-gov       960
## 3 Local-gov        2093
## 4 Never-worked        7
## 5 Private         22696
## 6 Self-emp-inc     1116
## 7 Self-emp-not-inc 2541
## 8 State-gov        1298
## 9 Without-pay        14
```

```
df %>% summarize(sum(!is.na(workclass)))
```

```
## # A tibble: 1 x 1
##   `sum(!is.na(workclass))`
##                      <int>
## 1                    32561
```

```
df %>% group_by(workclass) %>% summarize(first(age))
```

```
## # A tibble: 9 x 2
##   workclass       `first(age)`
##   <chr>                  <dbl>
## 1 ?                         54
## 2 Federal-gov               35
## 3 Local-gov                 56
## 4 Never-worked              18
## 5 Private                   38
## 6 Self-emp-inc              47
## 7 Self-emp-not-inc          50
## 8 State-gov                 39
## 9 Without-pay               65
```

```
df %>% group_by(workclass) %>% summarize(sd(capital_gain))
```

```
## # A tibble: 9 x 2
##   workclass       `sd(capital_gain)`
##   <chr>                        <dbl>
## 1 ?                            5147.
## 2 Federal-gov                  4102.
## 3 Local-gov                    5775.
## 4 Never-worked                    0
## 5 Private                      6424.
## 6 Self-emp-inc                17977.
## 7 Self-emp-not-inc            10986.
## 8 State-gov                    3778.
## 9 Without-pay                  1301.
```

```
df %>% group_by(workclass) %>% summarize(quantile(age, 0.5))
```

```
## # A tibble: 9 x 2
##   workclass       `quantile(age, 0.5)`
##   <chr>                          <dbl>
## 1 ?                                 35
## 2 Federal-gov                       43
## 3 Local-gov                         41
## 4 Never-worked                      18
## 5 Private                           35
## 6 Self-emp-inc                      45
## 7 Self-emp-not-inc                  44
## 8 State-gov                         39
## 9 Without-pay                       57
```

We can also apply the summary over selected columns.

```
df %>% select(1, 3, 5, 11, 12, 13) %>% summarize(across(everything(), mean))
```

```
## # A tibble: 1 x 6
##     age  fnlwgt education_num capital_gain capital_loss hours_per_week
##   <dbl>   <dbl>         <dbl>        <dbl>        <dbl>          <dbl>
## 1  38.6 189778.          10.1        1078.         87.3           40.4
```

4.7 Arranging

To sort the data set.

```
df %>% arrange(native_country)
```

```
## # A tibble: 32,561 x 15
##      age workclass fnlwgt education     education_num marital_status   occupation
##    <dbl> <chr>      <dbl> <chr>                 <dbl> <chr>            <chr>
##  1    40 Private   121772 Assoc-voc                11 Married-civ-spo~ Craft-rep~
##  2    31 Private    84154 Some-college             10 Married-civ-spo~ Sales
##  3    18 Private   226956 HS-grad                   9 Never-married    Other-ser~
##  4    32 ?         293936 7th-8th                   4 Married-spouse-~ ?
##  5    30 Private   117747 HS-grad                   9 Married-civ-spo~ Sales
##  6    56 Private   203580 HS-grad                   9 Married-civ-spo~ Adm-cleri~
##  7    45 Private   153141 HS-grad                   9 Married-civ-spo~ Adm-cleri~
##  8    39 ?         157443 Masters                  14 Married-civ-spo~ ?
##  9    34 State-gov  98101 Bachelors                13 Married-civ-spo~ Exec-mana~
## 10    42 Private   197583 Assoc-acdm               12 Married-civ-spo~ Exec-mana~
## # i 32,551 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

It is most useful to sort the data set after grouping and summarizing.

```
df %>%
  group_by(education) %>%
  summarize(count=n(),
            avg_net_gain = mean(capital_gain - capital_loss)) %>%
  arrange(desc(avg_net_gain))
```

```
## # A tibble: 16 x 3
##    education    count avg_net_gain
##    <chr>        <int>        <dbl>
##  1 Prof-school    576       10183.
##  2 Doctorate      413        4507.
##  3 Masters       1723        2396.
##  4 Bachelors     5355        1638.
##  5 Preschool       51         832.
##  6 Assoc-voc     1382         642.
##  7 Assoc-acdm    1067         547.
```

```
##  8 Some-college  7291        527.
##  9 HS-grad      10501        506.
## 10 10th           933        348.
## 11 9th            514        313.
## 12 12th           433        252.
## 13 7th-8th        646        168.
## 14 11th          1175        165.
## 15 5th-6th        333        108.
## 16 1st-4th        168         77.5
```

4.8 Separating and uniting

This is often useful to create new columns.

```
df %>% separate(target, into=c("sign", "amount"), sep="\\b")
```

```
## Warning: Expected 2 pieces. Additional pieces discarded in 32561 rows [1, 2, 3, 4, 5, 6,
## 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, ...].
```

```
## # A tibble: 32,561 x 16
##      age workclass       fnlwgt education education_num marital_status occupation
##    <dbl> <chr>            <dbl> <chr>             <dbl> <chr>          <chr>
##  1    39 State-gov        77516 Bachelors            13 Never-married  Adm-cleri~
##  2    50 Self-emp-not-~   83311 Bachelors            13 Married-civ-s~ Exec-mana~
##  3    38 Private         215646 HS-grad               9 Divorced       Handlers-~
##  4    53 Private         234721 11th                  7 Married-civ-s~ Handlers-~
##  5    28 Private         338409 Bachelors            13 Married-civ-s~ Prof-spec~
##  6    37 Private         284582 Masters              14 Married-civ-s~ Exec-mana~
##  7    49 Private         160187 9th                   5 Married-spous~ Other-ser~
##  8    52 Self-emp-not-~  209642 HS-grad               9 Married-civ-s~ Exec-mana~
##  9    31 Private          45781 Masters              14 Never-married  Prof-spec~
## 10    42 Private         159449 Bachelors            13 Married-civ-s~ Exec-mana~
## # i 32,551 more rows
## # i 9 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, sign <chr>, amount <chr>
```

```
df %>% unite(sex, race, age, col="description", sep="_", remove=FALSE)
```

```
## # A tibble: 32,561 x 16
##    description      age workclass fnlwgt education education_num marital_status
##    <chr>          <dbl> <chr>      <dbl> <chr>             <dbl> <chr>
##  1 Male_White_39     39 State-gov  77516 Bachelors            13 Never-married
##  2 Male_White_50     50 Self-emp~  83311 Bachelors            13 Married-civ-s~
##  3 Male_White_38     38 Private   215646 HS-grad               9 Divorced
##  4 Male_Black_53     53 Private   234721 11th                  7 Married-civ-s~
##  5 Female_Black_28   28 Private   338409 Bachelors            13 Married-civ-s~
##  6 Female_White_37   37 Private   284582 Masters              14 Married-civ-s~
##  7 Female_Black_49   49 Private   160187 9th                   5 Married-spous~
##  8 Male_White_52     52 Self-emp~ 209642 HS-grad               9 Married-civ-s~
##  9 Female_White_31   31 Private    45781 Masters              14 Never-married
## 10 Male_White_42     42 Private   159449 Bachelors            13 Married-civ-s~
```

```
## # i 32,551 more rows
## # i 9 more variables: occupation <chr>, relationship <chr>, race <chr>,
## #   sex <chr>, capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

4.8 Mutate function

Designed to create new variables.

```
df %>%
  mutate(total_gain = capital_gain - capital_loss,
         tax = ifelse(total_gain >= 15000,
                      total_gain * 0.1,
                      0)
  ) %>%
  arrange(desc(tax))
```

```
## # A tibble: 32,561 x 17
##      age workclass      fnlwgt education education_num marital_status occupation
##    <dbl> <chr>           <dbl> <chr>             <dbl> <chr>          <chr>
## 1     54 Self-emp-inc   166459 Prof-sch~            15 Married-civ-s~ Prof-spec~
## 2     52 Private        152234 HS-grad               9 Married-civ-s~ Exec-mana~
## 3     53 Self-emp-inc   263925 HS-grad               9 Married-civ-s~ Sales
## 4     52 Private        118025 Bachelors            13 Married-civ-s~ Exec-mana~
## 5     46 Private        370119 Prof-sch~            15 Married-civ-s~ Prof-spec~
## 6     43 Private        176270 Bachelors            13 Married-civ-s~ Exec-mana~
## 7     49 Private        159816 Bachelors            13 Married-civ-s~ Prof-spec~
## 8     50 Private        171338 Some-col~            10 Married-civ-s~ Exec-mana~
## 9     22 Self-emp-not-~ 202920 HS-grad               9 Never-married  Prof-spec~
## 10    43 Self-emp-inc   172826 Some-col~            10 Married-civ-s~ Sales
## # i 32,551 more rows
## # i 10 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>, total_gain <dbl>, tax <dbl>
```

We can use the mutate function to replace values (e.g. "?" by NA).

```
for (variable in colnames(df)) {
  print (
    paste (variable, dim(df[df[variable]=="?", variable])[1])
  )
}
```

```
## [1] "age 0"
## [1] "workclass 1836"
## [1] "fnlwgt 0"
## [1] "education 0"
## [1] "education_num 0"
## [1] "marital_status 0"
## [1] "occupation 1843"
## [1] "relationship 0"
## [1] "race 0"
```

```
## [1] "sex 0"
## [1] "capital_gain 0"
## [1] "capital_loss 0"
## [1] "hours_per_week 0"
## [1] "native_country 583"
## [1] "target 0"
```

```r
df_replaced <- df %>%
  mutate(workclass = replace(workclass, workclass=="?", NA),
         occupation = replace(occupation, occupation=="?", NA),
         native_country =replace(native_country, native_country=="?", NA)
  )
```

```r
for (variable in colnames(df_replaced)) {
  print (
    paste (variable, dim(df_replaced[df_replaced[variable]=="?", variable])[1])
  )
}
```

```
## [1] "age 0"
## [1] "workclass 1836"
## [1] "fnlwgt 0"
## [1] "education 0"
## [1] "education_num 0"
## [1] "marital_status 0"
## [1] "occupation 1843"
## [1] "relationship 0"
## [1] "race 0"
## [1] "sex 0"
## [1] "capital_gain 0"
## [1] "capital_loss 0"
## [1] "hours_per_week 0"
## [1] "native_country 583"
## [1] "target 0"
```

```r
df %>% mutate(workclass = na_if(workclass, "?"),
              occupation = na_if(occupation, "?"),
              native_country = na_if(native_country, "?"))
```

```
## # A tibble: 32,561 x 15
##       age workclass      fnlwgt education education_num marital_status occupation
##     <dbl> <chr>           <dbl> <chr>             <dbl> <chr>          <chr>
## 1      39 State-gov       77516 Bachelors            13 Never-married  Adm-cleri~
## 2      50 Self-emp-not-~  83311 Bachelors            13 Married-civ-s~ Exec-mana~
## 3      38 Private        215646 HS-grad               9 Divorced       Handlers-~
## 4      53 Private        234721 11th                  7 Married-civ-s~ Handlers-~
## 5      28 Private        338409 Bachelors            13 Married-civ-s~ Prof-spec~
## 6      37 Private        284582 Masters              14 Married-civ-s~ Exec-mana~
## 7      49 Private        160187 9th                   5 Married-spous~ Other-ser~
## 8      52 Self-emp-not-~ 209642 HS-grad               9 Married-civ-s~ Exec-mana~
## 9      31 Private         45781 Masters              14 Never-married  Prof-spec~
## 10     42 Private        159449 Bachelors            13 Married-civ-s~ Exec-mana~
```

```
## # i 32,551 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

Additional examples:

```
df %>%
  mutate(over_under = recode(target, "<=50K"="under",
                                      ">50K"="over")) %>%
  select(target, over_under)
```

```
## # A tibble: 32,561 x 2
##    target over_under
##    <chr>  <chr>
##  1 <=50K  under
##  2 <=50K  under
##  3 <=50K  under
##  4 <=50K  under
##  5 <=50K  under
##  6 <=50K  under
##  7 <=50K  under
##  8 >50K   over
##  9 >50K   over
## 10 >50K   over
## # i 32,551 more rows
```

```
df %>%
  mutate(age_avg = mean(age),
         over_under_age_avg = cut(age,
                                  c(0, mean(age), max(age)),
                                  c("Lower than avg", "Above the avg"))
  ) %>%
  select(age, age_avg, over_under_age_avg)
```

```
## # A tibble: 32,561 x 3
##      age age_avg over_under_age_avg
##    <dbl>   <dbl> <fct>
##  1    39    38.6 Above the avg
##  2    50    38.6 Above the avg
##  3    38    38.6 Lower than avg
##  4    53    38.6 Above the avg
##  5    28    38.6 Lower than avg
##  6    37    38.6 Lower than avg
##  7    49    38.6 Above the avg
##  8    52    38.6 Above the avg
##  9    31    38.6 Lower than avg
## 10    42    38.6 Above the avg
## # i 32,551 more rows
```