

# Tidyverse tutorial 1 - Basic operations

Ariane Ducellier

10/03/2023

Load R packages.

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.3      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.3      v tibble    3.2.1
## v lubridate  1.9.2      v tidyr     1.3.0
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(skimr)
```

## 1. Read data (csv format)

Read csv files with basic R.

```
df <- read.csv("../data/gapminder-data.csv")
print(class(df))
```

```
## [1] "data.frame"
```

Read csv files with tidyverse.

```
df_t <- read_csv("../data/gapminder-data.csv")
```

```
## New names:
## Rows: 1512 Columns: 10
## -- Column specification
## ----- Delimiter: "," chr
## (1): Country dbl (9): ...1, Year, gdp_per_capita,
## Electricity_consumption_per_capita, und...
## i Use 'spec()' to retrieve the full column specification for this data. i
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.
## * ' -> '...1'
```

```
print(class(df_t))
```

```
## [1] "spec_tbl_df" "tbl_df"      "tbl"        "data.frame"
```

```
df_t_sub <- read_csv("../data/gapminder-data.csv",
  col_select=c("Country", "Year", "gdp_per_capita"),
  na=c("", "NA"))
```

```
## New names:
## Rows: 1512 Columns: 3
## -- Column specification
## ----- Delimiter: "," chr
## (1): Country dbl (2): Year, gdp_per_capita
## i Use 'spec()' to retrieve the full column specification for this data. i
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.
## * ' -> '...1'
```

## 2. Get basic information on the data.

Reads the first part of a data frame.

```
df_t <- read_csv("../data/gapminder-data.csv")
```

```
## New names:
## Rows: 1512 Columns: 10
## -- Column specification
## ----- Delimiter: "," chr
## (1): Country dbl (9): ...1, Year, gdp_per_capita,
## Electricity_consumption_per_capita, und...
## i Use 'spec()' to retrieve the full column specification for this data. i
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.
## * ' -> '...1'
```

```
head(df_t, 3)
```

```
## # A tibble: 3 x 10
##   ...1 Country Year gdp_per_capita Electricity_consumption_p~1 under5mortality
##   <dbl> <chr>   <dbl>         <dbl>             <dbl>             <dbl>
## 1     0 Brazil  1800           1109                NA                417.
## 2     1 Brazil  1801           1109                NA                417.
## 3     2 Brazil  1802           1109                NA                417.
## # i abbreviated name: 1: Electricity_consumption_per_capita
## # i 4 more variables: Poverty <dbl>, BMI_male <dbl>, BMI_female <dbl>,
## #   population <dbl>
```

Reads the last part of a data frame.

```
tail(df_t, 3)
```

```
## # A tibble: 3 x 10
##   ...1 Country      Year gdp_per_capita Electricity_consumpt~1 under5mortality
##   <dbl> <chr>      <dbl>      <dbl>      <dbl>      <dbl>
## 1  1509 United Stat~ 2013      51282      NA          6.9
## 2  1510 United Stat~ 2014      52118      NA          6.7
## 3  1511 United Stat~ 2015      53354      NA          6.5
## # i abbreviated name: 1: Electricity_consumption_per_capita
## # i 4 more variables: Poverty <dbl>, BMI_male <dbl>, BMI_female <dbl>,
## #   population <dbl>
```

Gets column specifications of a tibble.

```
spec(df_t)
```

```
## cols(
##   ...1 = col_double(),
##   Country = col_character(),
##   Year = col_double(),
##   gdp_per_capita = col_double(),
##   Electricity_consumption_per_capita = col_double(),
##   under5mortality = col_double(),
##   Poverty = col_double(),
##   BMI_male = col_double(),
##   BMI_female = col_double(),
##   population = col_double()
## )
```

Prints the data: number of rows and columns, type of columns, and first rows.

```
glimpse(df_t)
```

```
## Rows: 1,512
## Columns: 10
## $ ...1      <dbl> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 1~
## $ Country   <chr> "Brazil", "Brazil", "Brazil", "Braz~
## $ Year       <dbl> 1800, 1801, 1802, 1803, 1804, 1805,~
## $ gdp_per_capita <dbl> 1109, 1109, 1109, 1109, 1109, 1110,~
## $ Electricity_consumption_per_capita <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ under5mortality <dbl> 417.44, 417.44, 417.44, 417.44, 417~
## $ Poverty    <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ BMI_male    <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ BMI_female  <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA,~
## $ population  <dbl> 3639636, NA, NA, NA, NA, NA, NA, NA, NA~
```

Returns descriptive statistics on each column of a data frame.

```
summary(df_t)
```

```
##      ...1      Country      Year      gdp_per_capita
## Min.   : 0.0      Length:1512      Min.   :1800      Min.   : 529
## 1st Qu.: 377.8    Class :character      1st Qu.:1854      1st Qu.: 1124
## Median : 755.5    Mode  :character      Median :1908      Median : 2496
## Mean   : 755.5                                Mean   :1908      Mean   : 7234
## 3rd Qu.:1133.2                                3rd Qu.:1961      3rd Qu.: 8219
## Max.   :1511.0                                Max.   :2015      Max.   :53354
##
## Electricity_consumption_per_capita under5mortality      Poverty
## Min.   : 97.78                                Min.   : 2.70      Min.   : 0.000
## 1st Qu.:1062.24                                1st Qu.: 77.59      1st Qu.: 0.920
## Median : 4310.62                                Median :306.66      Median : 9.385
## Mean   : 4386.74                                Mean   :260.02      Mean   :15.338
## 3rd Qu.: 6495.64                                3rd Qu.:417.44      3rd Qu.:15.960
## Max.   :13704.58                                Max.   :539.16      Max.   :84.270
## NA's   :1181                                NA's   :1440
## BMI_male      BMI_female      population
## Min.   :20.62      Min.   :20.48      Min.   :3.640e+06
## 1st Qu.:22.22      1st Qu.:21.90      1st Qu.:6.740e+07
## Median :24.04      Median :24.57      Median :1.250e+08
## Mean   :24.16      Mean   :23.91      Mean   :2.996e+08
## 3rd Qu.:26.17      3rd Qu.:25.56      3rd Qu.:3.767e+08
## Max.   :28.46      Max.   :28.34      Max.   :1.376e+09
## NA's   :1309      NA's   :1309      NA's   :945
```

Provides a broad overview of a data frame, handles data of all types, dispatching a different set of summary functions based on the types of columns in the data frame.

```
#skim(df_t)
```

### 3. The pipe operator

```
data(iris)
df_iris <- iris %>%
  group_by(Species) %>%
  summarize_if(is.numeric, mean) %>%
  ungroup() %>%
  gather(measure, value, -Species) %>%
  arrange(value)
```

```
data(iris)
df_iris_alt <- group_by(iris, Species)
df_iris_alt <- summarize_if(df_iris_alt, is.numeric, mean)
df_iris_alt <- ungroup(df_iris_alt)
df_iris_alt <- gather(df_iris_alt, measure, value, -Species)
df_iris_alt <- arrange(df_iris_alt, value)
```

## 4. Transform the data

```
header <- c("age", "workclass", "fnlwgt", "education",
  "education_num", "marital_status", "occupation",
  "relationship", "race", "sex", "capital_gain",
  "capital_loss", "hours_per_week", "native_country", "target")
df <- read_csv("https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data",
  col_names=header, trim_ws=TRUE)
```

```
## Rows: 32561 Columns: 15
## -- Column specification -----
## Delimiter: ","
## chr (9): workclass, education, marital_status, occupation, relationship, rac...
## dbl (6): age, fnlwgt, education_num, capital_gain, capital_loss, hours_per_week
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

### 4.1 Slicing

Cuts unwanted parts of the data set.

```
df %>% slice_head(n=5)
```

```
## # A tibble: 5 x 15
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>          <dbl> <chr>          <dbl> <chr>          <chr>
## 1   39 State-gov      77516 Bachelors          13 Never-married Adm-cleri-
## 2   50 Self-emp-not-i~ 83311 Bachelors          13 Married-civ-s~ Exec-mana~
## 3   38 Private      215646 HS-grad           9 Divorced      Handlers--
## 4   53 Private      234721 11th             7 Married-civ-s~ Handlers--
## 5   28 Private      338409 Bachelors          13 Married-civ-s~ Prof-spec~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_tail(n=5)
```

```
## # A tibble: 5 x 15
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>          <dbl> <chr>          <dbl> <chr>          <chr>
## 1   27 Private      257302 Assoc-acdm          12 Married-civ-spo~ Tech-sup~
## 2   40 Private      154374 HS-grad           9 Married-civ-spo~ Machine-o~
## 3   58 Private      151910 HS-grad           9 Widowed        Adm-cleri-
## 4   22 Private      201490 HS-grad           9 Never-married Adm-cleri-
## 5   52 Self-emp-inc 287927 HS-grad           9 Married-civ-spo~ Exec-mana~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_min(age, prop=0.10)
```

```
## # A tibble: 3,895 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>    <dbl> <chr>          <dbl> <chr>          <chr>
## 1  17 ?          304873 10th              6 Never-married ?
## 2  17 Private    65368 11th              7 Never-married Sales
## 3  17 Private   245918 11th              7 Never-married Other-service
## 4  17 Private   191260 9th               5 Never-married Other-service
## 5  17 Private   270942 5th-6th           3 Never-married Other-service
## 6  17 Private    89821 11th              7 Never-married Other-service
## 7  17 Private   175024 11th              7 Never-married Handlers-clean~
## 8  17 ?          202521 11th              7 Never-married ?
## 9  17 ?          258872 11th              7 Never-married ?
## 10 17 Private    211870 9th               5 Never-married Other-service
## # i 3,885 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_max(age, prop=0.30)
```

```
## # A tibble: 10,361 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>    <dbl> <chr>          <dbl> <chr>          <chr>
## 1  90 Private    51744 HS-grad           9 Never-married Other-ser~
## 2  90 Private   137018 HS-grad           9 Never-married Other-ser~
## 3  90 Private   221832 Bachelors        13 Married-civ-spo~ Exec-mana~
## 4  90 Private    52386 Some-college  10 Never-married Other-ser~
## 5  90 Private   171956 Some-college  10 Separated      Adm-cleri~
## 6  90 Private   313986 11th              7 Never-married Handlers--
## 7  90 ?          256514 Bachelors        13 Widowed        ?
## 8  90 Private    52386 Some-college  10 Never-married Other-ser~
## 9  90 Private   141758 9th               5 Never-married Adm-cleri~
## 10 90 Local-gov  227796 Masters        14 Married-civ-spo~ Exec-mana~
## # i 10,351 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

```
df %>% slice_sample(n=10, replace=TRUE)
```

```
## # A tibble: 10 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>    <dbl> <chr>          <dbl> <chr>          <chr>
## 1  51 Private    57698 HS-grad           9 Married-spous~ Other-ser~
## 2  47 Private   168232 Bachelors        13 Married-civ-s~ Adm-cleri~
## 3  27 Private   111900 Some-colle~  10 Never-married Prof-spec~
## 4  35 Private   187625 Some-colle~  10 Never-married Other-ser~
## 5  26 Private   132572 Bachelors        13 Never-married Adm-cleri~
## 6  43 Self-emp-inc 188436 Masters        14 Married-civ-s~ Exec-mana~
```

```
## 7 35 Private 241998 Bachelors 13 Married-civ-s~ Exec-mana~
## 8 41 Private 171546 Bachelors 13 Married-civ-s~ Tech-supp~
## 9 19 Private 118352 Some-colle~ 10 Never-married Other-ser~
## 10 33 Private 176992 Masters 14 Married-civ-s~ Prof-spec~
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

## 4.2 Filtering

Apply a condition to one of the variables to filter unwanted rows of the data.

```
df %>% filter(age > 30)
```

```
## # A tibble: 21,989 x 15
##   age workclass   fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>          <dbl> <chr>         <chr>
## 1 39 State-gov    77516 Bachelors        13 Never-married Adm-cleri~
## 2 50 Self-emp-not~ 83311 Bachelors        13 Married-civ-s~ Exec-mana~
## 3 38 Private    215646 HS-grad         9 Divorced      Handlers~
## 4 53 Private    234721 11th            7 Married-civ-s~ Handlers~
## 5 37 Private    284582 Masters        14 Married-civ-s~ Exec-mana~
## 6 49 Private    160187 9th             5 Married-spous~ Other-ser~
## 7 52 Self-emp-not~ 209642 HS-grad         9 Married-civ-s~ Exec-mana~
## 8 31 Private     45781 Masters        14 Never-married Prof-spec~
## 9 42 Private    159449 Bachelors        13 Married-civ-s~ Exec-mana~
## 10 37 Private    280464 Some-col~    10 Married-civ-s~ Exec-mana~
## # i 21,979 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## # capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## # native_country <chr>, target <chr>
```

## 4.3 Selecting

Select specific columns.

```
df %>% select(marital_status, age)
```

```
## # A tibble: 32,561 x 2
##   marital_status   age
##   <chr>         <dbl>
## 1 Never-married    39
## 2 Married-civ-spouse 50
## 3 Divorced        38
## 4 Married-civ-spouse 53
## 5 Married-civ-spouse 28
## 6 Married-civ-spouse 37
## 7 Married-spouse-absent 49
## 8 Married-civ-spouse 52
## 9 Never-married    31
## 10 Married-civ-spouse 42
## # i 32,551 more rows
```

## 4.4 Unique values

Get unique entries for categorical variables.

```
df %>% distinct(sex)
```

```
## # A tibble: 2 x 1
##   sex
##   <chr>
## 1 Male
## 2 Female
```

## 4.5 Grouping

Group by column and summarize.

```
df %>%
  group_by(workclass) %>%
  summarize(age_avg=mean(age))
```

```
## # A tibble: 9 x 2
##   workclass      age_avg
##   <chr>         <dbl>
## 1 ?             41.0
## 2 Federal-gov   42.6
## 3 Local-gov     41.8
## 4 Never-worked  20.6
## 5 Private       36.8
## 6 Self-emp-inc  46.0
## 7 Self-emp-not-inc 45.0
## 8 State-gov     39.4
## 9 Without-pay   47.8
```

## 4.6 Summarizing

The summary may be: - counting observations - counting available observations (i.e. not NA) - getting first or last value - compute statistics on each group (mean, standard deviation, quantile)

```
df %>% group_by(workclass) %>% summarize(n())
```

```
## # A tibble: 9 x 2
##   workclass      'n()'
##   <chr>         <int>
## 1 ?             1836
## 2 Federal-gov   960
## 3 Local-gov     2093
## 4 Never-worked    7
## 5 Private      22696
## 6 Self-emp-inc   1116
## 7 Self-emp-not-inc 2541
## 8 State-gov     1298
## 9 Without-pay    14
```



```
df %>% summarize(sum(!is.na(workclass)))
```

```
## # A tibble: 1 x 1
##   'sum(!is.na(workclass))'
##   <int>
## 1      32561
```

```
df %>% group_by(workclass) %>% summarize(first(age))
```

```
## # A tibble: 9 x 2
##   workclass      'first(age)'
##   <chr>          <dbl>
## 1 ?              54
## 2 Federal-gov    35
## 3 Local-gov      56
## 4 Never-worked   18
## 5 Private        38
## 6 Self-emp-inc    47
## 7 Self-emp-not-inc 50
## 8 State-gov      39
## 9 Without-pay    65
```

```
df %>% group_by(workclass) %>% summarize(sd(capital_gain))
```

```
## # A tibble: 9 x 2
##   workclass      'sd(capital_gain)'
##   <chr>          <dbl>
## 1 ?              5147.
## 2 Federal-gov    4102.
## 3 Local-gov      5775.
## 4 Never-worked   0
## 5 Private        6424.
## 6 Self-emp-inc    17977.
## 7 Self-emp-not-inc 10986.
## 8 State-gov      3778.
## 9 Without-pay    1301.
```

```
df %>% group_by(workclass) %>% summarize(quantile(age, 0.5))
```

```
## # A tibble: 9 x 2
##   workclass      'quantile(age, 0.5)'
##   <chr>          <dbl>
## 1 ?              35
## 2 Federal-gov    43
## 3 Local-gov      41
## 4 Never-worked   18
## 5 Private        35
## 6 Self-emp-inc    45
## 7 Self-emp-not-inc 44
## 8 State-gov      39
## 9 Without-pay    57
```

We can also apply the summary over selected columns.

```
df %>% select(1, 3, 5, 11, 12, 13) %>% summarize(across(everything(), mean))
```

```
## # A tibble: 1 x 6
##   age  fnlwgt education_num capital_gain capital_loss hours_per_week
##   <dbl> <dbl>         <dbl>         <dbl>         <dbl>         <dbl>
## 1  38.6 189778.         10.1         1078.         87.3         40.4
```

## 4.7 Arranging

To sort the data set.

```
df %>% arrange(native_country)
```

```
## # A tibble: 32,561 x 15
##   age workclass fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>         <dbl> <chr>         <chr>
## 1  40 Private   121772 Assoc-voc         11 Married-civ-spo~ Craft-rep~
## 2  31 Private   84154 Some-college    10 Married-civ-spo~ Sales
## 3  18 Private  226956 HS-grad          9 Never-married   Other-ser~
## 4  32 ?        293936 7th-8th          4 Married-spouse~ ?
## 5  30 Private  117747 HS-grad          9 Married-civ-spo~ Sales
## 6  56 Private  203580 HS-grad          9 Married-civ-spo~ Adm-cleri~
## 7  45 Private  153141 HS-grad          9 Married-civ-spo~ Adm-cleri~
## 8  39 ?        157443 Masters        14 Married-civ-spo~ ?
## 9  34 State-gov 98101 Bachelors      13 Married-civ-spo~ Exec-mana~
## 10 42 Private  197583 Assoc-acdm      12 Married-civ-spo~ Exec-mana~
## # i 32,551 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

It is most useful to sort the data set after grouping and summarizing.

```
df %>%
  group_by(education) %>%
  summarize(count=n(),
            avg_net_gain = mean(capital_gain - capital_loss)) %>%
  arrange(desc(avg_net_gain))
```

```
## # A tibble: 16 x 3
##   education count avg_net_gain
##   <chr>      <int>      <dbl>
## 1 Prof-school  576    10183.
## 2 Doctorate   413    4507.
## 3 Masters    1723    2396.
## 4 Bachelors  5355    1638.
## 5 Preschool   51     832.
## 6 Assoc-voc  1382    642.
## 7 Assoc-acdm 1067    547.
```

```
## 8 Some-college 7291 527.
## 9 HS-grad 10501 506.
## 10 10th 933 348.
## 11 9th 514 313.
## 12 12th 433 252.
## 13 7th-8th 646 168.
## 14 11th 1175 165.
## 15 5th-6th 333 108.
## 16 1st-4th 168 77.5
```

## 4.8 Separating and uniting

This is often useful to create new columns.

```
df %>% separate(target, into=c("sign", "amount"), sep="\\b")
```

```
## Warning: Expected 2 pieces. Additional pieces discarded in 32561 rows [1, 2, 3, 4, 5, 6,
## 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, ...].
```

```
## # A tibble: 32,561 x 16
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>      <dbl> <chr>          <dbl> <chr>      <chr>
## 1 39 State-gov      77516 Bachelors         13 Never-married Adm-cleri~
## 2 50 Self-emp-not~ 83311 Bachelors         13 Married-civ-s~ Exec-mana~
## 3 38 Private      215646 HS-grad           9 Divorced      Handlers~
## 4 53 Private      234721 11th             7 Married-civ-s~ Handlers~
## 5 28 Private      338409 Bachelors         13 Married-civ-s~ Prof-spec~
## 6 37 Private      284582 Masters          14 Married-civ-s~ Exec-mana~
## 7 49 Private      160187 9th              5 Married-spous~ Other-ser~
## 8 52 Self-emp-not~ 209642 HS-grad           9 Married-civ-s~ Exec-mana~
## 9 31 Private      45781 Masters          14 Never-married Prof-spec~
## 10 42 Private     159449 Bachelors          13 Married-civ-s~ Exec-mana~
## # i 32,551 more rows
## # i 9 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, sign <chr>, amount <chr>
```

```
df %>% unite(sex, race, age, col="description", sep="_", remove=FALSE)
```

```
## # A tibble: 32,561 x 16
##   description      age workclass fnlwgt education education_num marital_status
##   <chr>      <dbl> <chr>      <dbl> <chr>          <dbl> <chr>
## 1 Male_White_39 39 State-gov 77516 Bachelors         13 Never-married
## 2 Male_White_50 50 Self-emp~ 83311 Bachelors         13 Married-civ-s~
## 3 Male_White_38 38 Private 215646 HS-grad           9 Divorced
## 4 Male_Black_53 53 Private 234721 11th             7 Married-civ-s~
## 5 Female_Black_28 28 Private 338409 Bachelors         13 Married-civ-s~
## 6 Female_White_37 37 Private 284582 Masters          14 Married-civ-s~
## 7 Female_Black_49 49 Private 160187 9th              5 Married-spous~
## 8 Male_White_52 52 Self-emp~ 209642 HS-grad           9 Married-civ-s~
## 9 Female_White_31 31 Private 45781 Masters          14 Never-married
## 10 Male_White_42 42 Private 159449 Bachelors          13 Married-civ-s~
```

```
## # i 32,551 more rows
## # i 9 more variables: occupation <chr>, relationship <chr>, race <chr>,
## #   sex <chr>, capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

## 4.8 Mutate function

Designed to create new variables.

```
df %>%
  mutate(total_gain = capital_gain - capital_loss,
         tax = ifelse(total_gain >= 15000,
                      total_gain * 0.1,
                      0)
  ) %>%
  arrange(desc(tax))
```

```
## # A tibble: 32,561 x 17
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>          <dbl> <chr>         <chr>
## 1    54 Self-emp-inc  166459 Prof-sch~      15 Married-civ-s~ Prof-spec~
## 2    52 Private      152234 HS-grad         9 Married-civ-s~ Exec-mana~
## 3    53 Self-emp-inc  263925 HS-grad         9 Married-civ-s~ Sales
## 4    52 Private      118025 Bachelors      13 Married-civ-s~ Exec-mana~
## 5    46 Private      370119 Prof-sch~      15 Married-civ-s~ Prof-spec~
## 6    43 Private      176270 Bachelors      13 Married-civ-s~ Exec-mana~
## 7    49 Private      159816 Bachelors      13 Married-civ-s~ Prof-spec~
## 8    50 Private      171338 Some-col~     10 Married-civ-s~ Exec-mana~
## 9    22 Self-emp-not-- 202920 HS-grad         9 Never-married Prof-spec~
## 10   43 Self-emp-inc  172826 Some-col~     10 Married-civ-s~ Sales
## # i 32,551 more rows
## # i 10 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>, total_gain <dbl>, tax <dbl>
```

We can use the mutate function to replace values (e.g. “?” by NA).

```
for (variable in colnames(df)) {
  print (
    paste (variable, dim(df[df[variable]=="?", variable])[1])
  )
}
```

```
## [1] "age 0"
## [1] "workclass 1836"
## [1] "fnlwgt 0"
## [1] "education 0"
## [1] "education_num 0"
## [1] "marital_status 0"
## [1] "occupation 1843"
## [1] "relationship 0"
## [1] "race 0"
```

```
## [1] "sex 0"
## [1] "capital_gain 0"
## [1] "capital_loss 0"
## [1] "hours_per_week 0"
## [1] "native_country 583"
## [1] "target 0"
```

```
df_replaced <- df %>%
  mutate(workclass = replace(workclass, workclass=="?", NA),
         occupation = replace(occupation, occupation=="?", NA),
         native_country = replace(native_country, native_country=="?", NA)
  )
```

```
for (variable in colnames(df_replaced)) {
  print (
    paste (variable, dim(df_replaced[df_replaced[variable]=="?", variable])[1])
  )
}
```

```
## [1] "age 0"
## [1] "workclass 1836"
## [1] "fnlwgt 0"
## [1] "education 0"
## [1] "education_num 0"
## [1] "marital_status 0"
## [1] "occupation 1843"
## [1] "relationship 0"
## [1] "race 0"
## [1] "sex 0"
## [1] "capital_gain 0"
## [1] "capital_loss 0"
## [1] "hours_per_week 0"
## [1] "native_country 583"
## [1] "target 0"
```

```
df %>% mutate(workclass = na_if(workclass, "?"),
              occupation = na_if(occupation, "?"),
              native_country = na_if(native_country, "?"))
```

```
## # A tibble: 32,561 x 15
##   age workclass      fnlwgt education education_num marital_status occupation
##   <dbl> <chr>         <dbl> <chr>          <dbl> <chr>         <chr>
## 1    39 State-gov      77516 Bachelors         13 Never-married Adm-cleri~
## 2    50 Self-emp-not-~ 83311 Bachelors         13 Married-civ-s~ Exec-mana~
## 3    38 Private       215646 HS-grad           9 Divorced      Handlers--
## 4    53 Private       234721 11th             7 Married-civ-s~ Handlers--
## 5    28 Private       338409 Bachelors         13 Married-civ-s~ Prof-spec~
## 6    37 Private       284582 Masters          14 Married-civ-s~ Exec-mana~
## 7    49 Private       160187 9th              5 Married-spous~ Other-ser~
## 8    52 Self-emp-not-~ 209642 HS-grad           9 Married-civ-s~ Exec-mana~
## 9    31 Private       45781 Masters          14 Never-married Prof-spec~
## 10   42 Private       159449 Bachelors         13 Married-civ-s~ Exec-mana~
```

```
## # i 32,551 more rows
## # i 8 more variables: relationship <chr>, race <chr>, sex <chr>,
## #   capital_gain <dbl>, capital_loss <dbl>, hours_per_week <dbl>,
## #   native_country <chr>, target <chr>
```

Additional examples:

```
df %>%
  mutate(over_under = recode(target, "<=50K"="under",
                                ">50K"="over")) %>%
  select(target, over_under)
```

```
## # A tibble: 32,561 x 2
##   target over_under
##   <chr>   <chr>
## 1 <=50K   under
## 2 <=50K   under
## 3 <=50K   under
## 4 <=50K   under
## 5 <=50K   under
## 6 <=50K   under
## 7 <=50K   under
## 8 >50K    over
## 9 >50K    over
## 10 >50K   over
## # i 32,551 more rows
```

```
df %>%
  mutate(age_avg = mean(age),
         over_under_age_avg = cut(age,
                                   c(0, mean(age), max(age)),
                                   c("Lower than avg", "Above the avg"))
  ) %>%
  select(age, age_avg, over_under_age_avg)
```

```
## # A tibble: 32,561 x 3
##   age age_avg over_under_age_avg
##   <dbl>   <dbl> <fct>
## 1   39   38.6 Above the avg
## 2   50   38.6 Above the avg
## 3   38   38.6 Lower than avg
## 4   53   38.6 Above the avg
## 5   28   38.6 Lower than avg
## 6   37   38.6 Lower than avg
## 7   49   38.6 Above the avg
## 8   52   38.6 Above the avg
## 9   31   38.6 Lower than avg
## 10  42   38.6 Above the avg
## # i 32,551 more rows
```

## 4.9 Joining tibbles

```
sales <- data.frame(
  date = c("2022-01-01", "2022-01-02", "2022-01-03", "2022-01-04", "2022-01-05"),
  store_cd= c(1, 2, 3, 4, 5),
  product_cd= c(1, 2, 3, 4, 5),
  qty= c(10, 12, 9, 12, 8),
  sales= c(30, 60, 45, 24, 32)
)

stores <- data.frame(
  store_cd= c(1, 2, 3, 4, 6),
  address= c("1 main st", "20 side st", "19 square blvd", "101 first st", "1002 retail ave"),
  city= c("Main", "East", "West", "North", "South"),
  open_hours= c("7-23", "7-23", "9-21", "9-21", "9-21")
)

products <- data.frame(
  product_cd= c(1, 2, 3, 4, 6),
  description= c("Soft drink", "Frozen snack", "Fruit", "Water", "Fruit 2"),
  unit_price= c(3.0, 5.0, 5.0, 2.0, 4.0),
  unit_measure= c("each", "each", "kg", "each", "kg")
)
```

### 4.9.1 Left join

All the rows from sales and matched rows from products.

```
sales %>% left_join(products, by="product_cd")
```

##	date	store_cd	product_cd	qty	sales	description	unit_price	unit_measure
## 1	2022-01-01	1	1	10	30	Soft drink	3	each
## 2	2022-01-02	2	2	12	60	Frozen snack	5	each
## 3	2022-01-03	3	3	9	45	Fruit	5	kg
## 4	2022-01-04	4	4	12	24	Water	2	each
## 5	2022-01-05	5	5	8	32	<NA>	NA	<NA>

### 4.9.2 Right join

All the rows from stores and matched rows from sales.

```
sales %>% right_join(stores, by="store_cd")
```

##	date	store_cd	product_cd	qty	sales	address	city	open_hours
## 1	2022-01-01	1	1	10	30	1 main st	Main	7-23
## 2	2022-01-02	2	2	12	60	20 side st	East	7-23
## 3	2022-01-03	3	3	9	45	19 square blvd	West	9-21
## 4	2022-01-04	4	4	12	24	101 first st	North	9-21
## 5	<NA>	6	NA	NA	NA	1002 retail ave	South	9-21

### 4.9.3 Inner join

All the rows common to sales and stores.

```
sales %>% inner_join(stores, by="store_cd")
```

```
##           date store_cd product_cd qty sales           address city open_hours
## 1 2022-01-01         1           1  10    30         1 main st  Main       7-23
## 2 2022-01-02         2           2  12    60        20 side st  East       7-23
## 3 2022-01-03         3           3   9    45 19 square blvd  West       9-21
## 4 2022-01-04         4           4  12    24 101 first st  North      9-21
```

### 4.9.4 Full join

All the rows from sales and stores.

```
sales %>% full_join(stores)
```

```
## Joining with 'by = join_by(store_cd)'
```

```
##           date store_cd product_cd qty sales           address city open_hours
## 1 2022-01-01         1           1  10    30         1 main st  Main       7-23
## 2 2022-01-02         2           2  12    60        20 side st  East       7-23
## 3 2022-01-03         3           3   9    45 19 square blvd  West       9-21
## 4 2022-01-04         4           4  12    24 101 first st  North      9-21
## 5 2022-01-05         5           5   8    32                <NA> <NA>      <NA>
## 6      <NA>         6          NA   NA   NA 1002 retail ave  South      9-21
```

### 4.9.5 Anti-join

Only rows that are in sales but not in products.

```
sales %>% anti_join(products)
```

```
## Joining with 'by = join_by(product_cd)'
```

```
##           date store_cd product_cd qty sales
## 1 2022-01-05         5           5   8    32
```

## 4.10 Reshaping tables

```
df_wide <- data.frame(
  project = c("project1", "project2", "project3"),
  Jan= sample(1000:2000, 3),
  Feb= sample(1000:2000, 3),
  Mar= sample(1000:2000, 3)
)
```

This is not a tidy data set (there is more than 1 observation per row).



```
df_long <- df_wide %>%  
  pivot_longer(cols= 2:4,  
               names_to = "months",  
               values_to = "expenses")
```

If we need a smaller table for visualization or for a presentation.

```
df_wide_2 <- df_long %>%  
  pivot_wider(names_from = "months",  
              values_from = "expenses")
```